

ELEC 424/553

Mobile & Embedded Systems

Lecture 4 - Process Management, Getting Comfortable with Linux, & Exploring Linux Source Code

Project 1: Cracking Open A Cold Raspberry Pi

- To be posted soon
- Setting up your RPi Zero W

CLEAR Storage

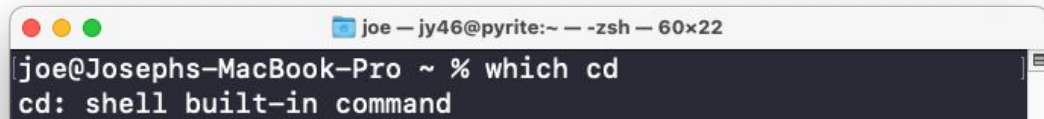
Everyone's CLEAR storage should have been increased by 2GB

IT WILL GO BACK DOWN BY 2GB AROUND DECEMBER 21st!

YOU COULD LOSE THINGS IF YOU DON'T DELETE THINGS BY THEN

Process: Program Instance

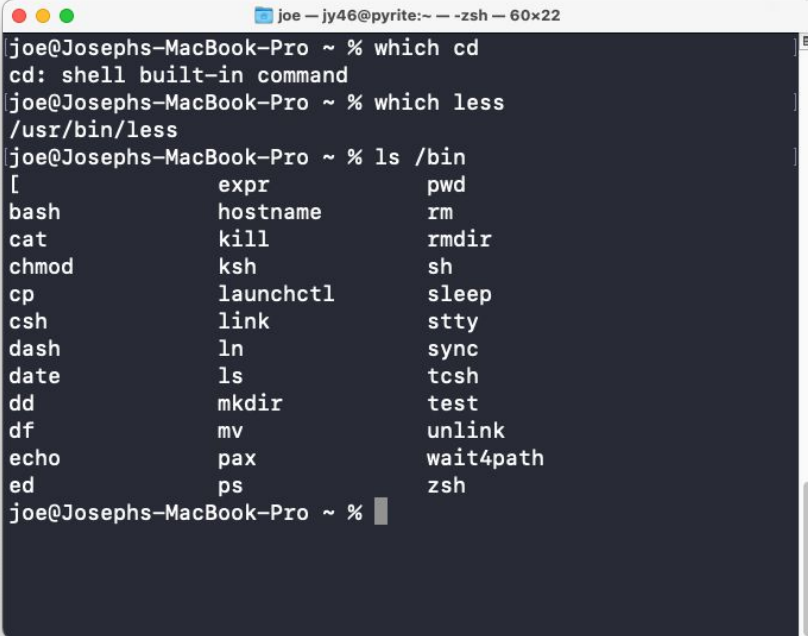
- Examples of programs:
 - Internal shell command
 - `cd`



```
joe — jy46@pyrite:~ — zsh — 60x22
joe@Josephs-MacBook-Pro ~ % which cd
cd: shell built-in command
```

Process: Program Instance

- Examples of programs:
 - Internal shell command
 - `cd`
 - Binary executable
 - `01001111101...`



```
joe — jy46@pyrite:~ — zsh — 60x22
joe@Josephs-MacBook-Pro ~ % which cd
cd: shell built-in command
joe@Josephs-MacBook-Pro ~ % which less
/usr/bin/less
joe@Josephs-MacBook-Pro ~ % ls /bin
[          expr          pwd
bash      hostname      rm
cat       kill          rmdir
chmod     ksh            sh
cp        launchctl     sleep
csh       link          stty
dash      ln            sync
date      ls             tcsh
dd        mkdir        test
df        mv           unlink
echo      pax           wait4path
ed        ps            zsh
joe@Josephs-MacBook-Pro ~ %
```

Process: Program Instance

- Examples of programs:

- Internal shell command

- `cd`

- Binary executable

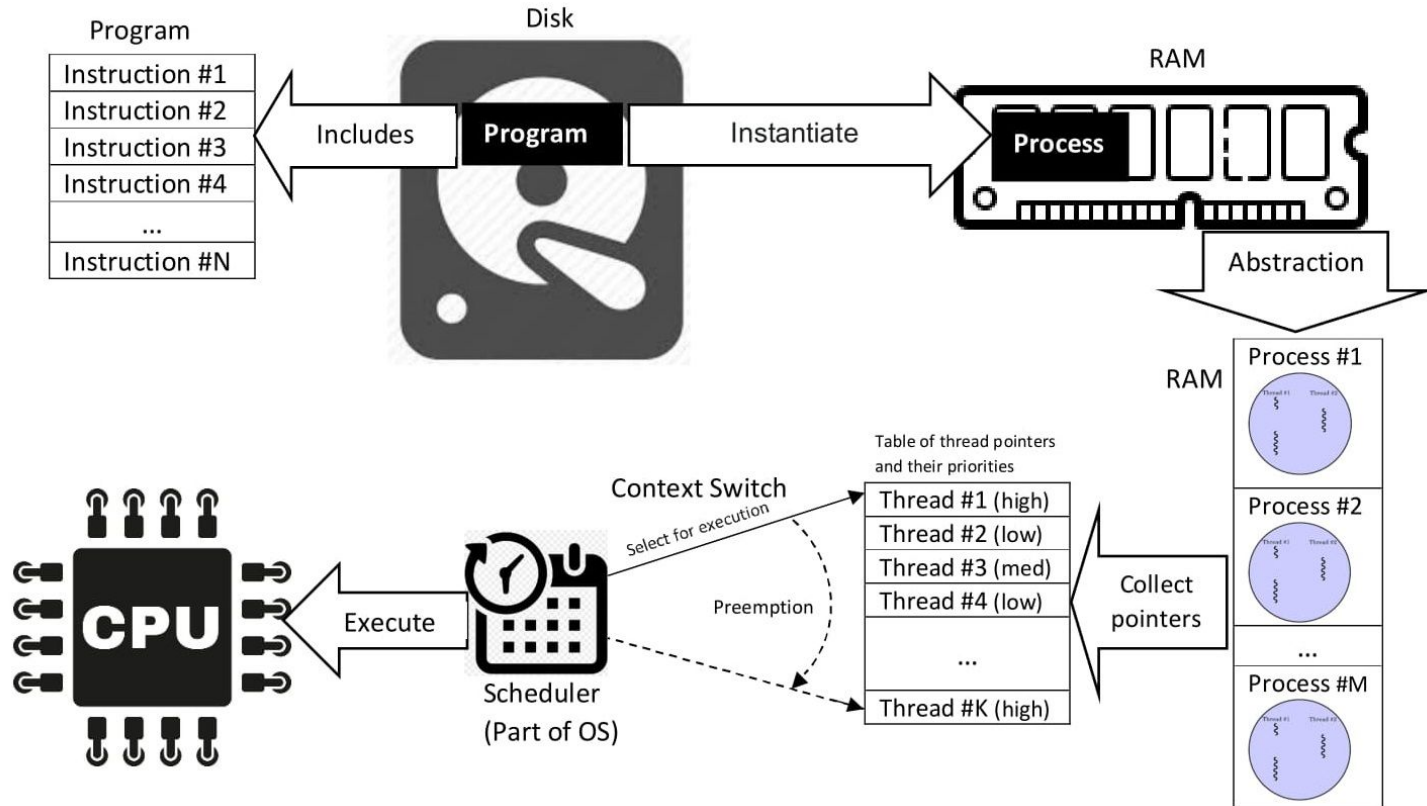
- `01001111101...`

- Shell script

- File of commands


https://linuxhint.com/30_bash_script_examples/

Programs, Processes, Threads, ...



Processes in Linux

- The goal of an OS is to host applications
- Process management is critical
- From Love (the author):
“A **process** is a program (object code stored on some media) in the midst of execution”
- Processes don't (generally) share memory

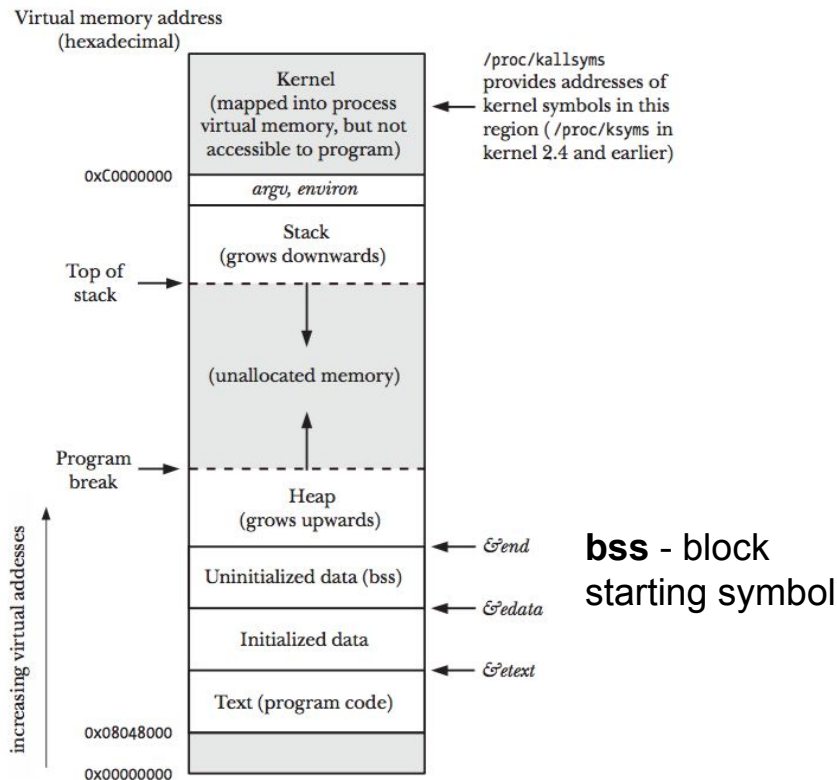


Author: Bobbo. CC Attribution-Share Alike 3.0 Unported license ([link](https://commons.wikimedia.org/wiki/File:Kernel_Layout.svg)).
Unmodified. URL:
https://commons.wikimedia.org/wiki/File:Kernel_Layout.svg

Process Virtualization

- Virtual processor
 - Makes it look like process has processor all to itself
 - Threads view private virtualized processors
- Virtual memory
 - Makes it look like process has total memory to itself
 - Threads view the same virtual memory

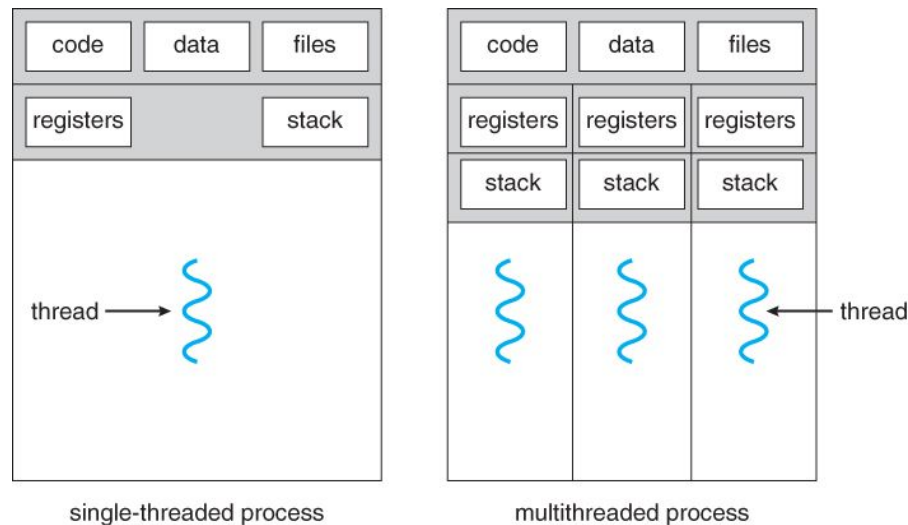
What Does a Process Consist Of?



- Not only the program code (“text”)
- Also contains:
 - Open files
 - Pending signals
 - Internal kernel data
 - Memory address space
 - 1+ threads
 - Data section

Threads

- **Threads** are *technically* the minimum set of instructions that a scheduler can handle
- Schedulers generally manage threads, but in Linux the situation is unique
- A thread gets its own:
 - Program counter
 - Stack
 - Registers
- **Multithreaded** programs have 2+ threads
- In contrast to processes, threads do share memory




URL: https://www.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/4_Threads.html

Figure source: Abraham Silberschatz, Greg Gagne, and Peter Baer Galvin, "Operating System Concepts, Ninth Edition ", Chapter 4

Multithreading in Software & Hardware

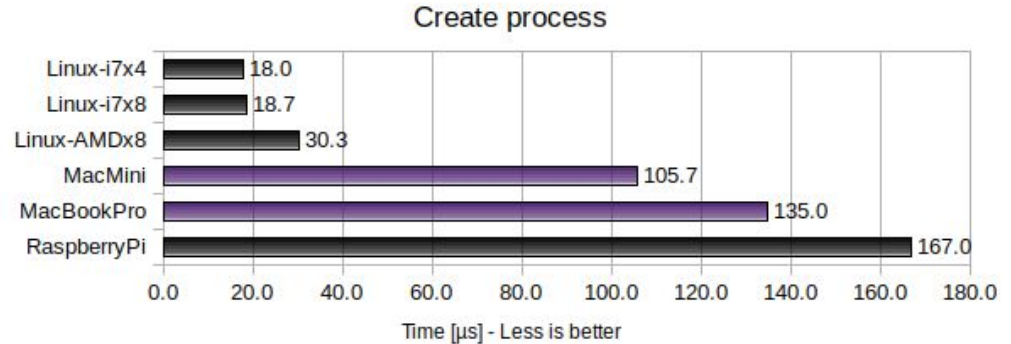
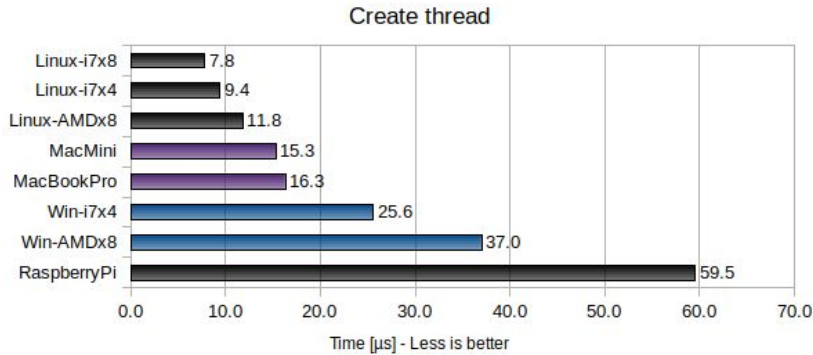
- It's possible to execute multiple threads at the same time
- Multithreading
 - Requires SW & HW support
- Multicore



Threads vs. Processes in Linux

- Linux views threads as particular processes
 - Processes with items common to other processes
 - Get their own `task_struct`
- Naming confusion: **Task**
 - Love says task and process are synonyms
 - Wikipedia will tell you task and thread are synonyms ([link](#))
- Windows different - threads particularly handled by kernel

Linux Thread/Process Creation Time




Running a Program (Creating a Process)

- nano hello.c

```
#include <stdio.h>
int main()
{
    printf("Hello, World!\n");
}
```

- gcc hello.c -o hello
- objdump -d hello
 - Assembly language
 - It's just instructions! :)

Viewing Existing Processes

- `top` shell command shows processes
- `shift+f` provides interactive menu
- Can press `s` on PID then `q`
- Then can do `shift R` to change to ascending
- Notice `systemd` 

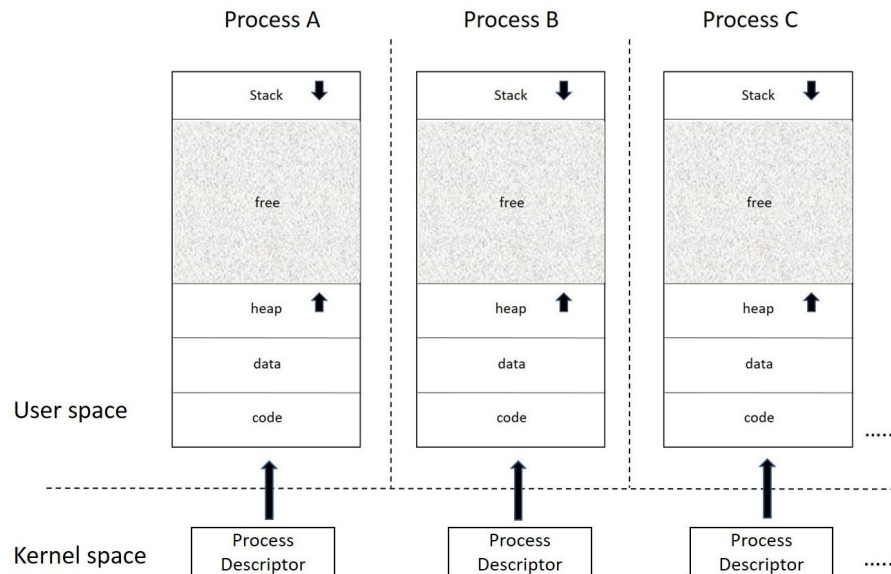
```
joe@joe-pink: ~  
(base) joe@joe-pink:~$ top
```

```
top - 12:41:47 up 2:06, 1 user, load average: 1.19, 0.88, 0.65  
Tasks: 462 total, 2 running, 460 sleeping, 0 stopped, 0 zombie  
%Cpu(s): 2.5 us, 1.1 sy, 0.0 ni, 96.4 id, 0.0 wa, 0.0 hi, 0.1 si, 0.0 st  
MiB Mem : 64212.4 total, 54843.7 free, 6039.1 used, 3329.7 buff/cache  
MiB Swap: 881.5 total, 881.5 free, 0.0 used, 57296.9 avail Mem
```

| PID | USER | PR | NI | VIRT | RES | SHR | S | %CPU | %MEM | TIME+ | COMMAND |
|-------|------|-----|-----|---------|--------|--------|---|------|------|----------|-----------------|
| 2533 | joe | 20 | 0 | 13.7g | 1.1g | 405904 | S | 29.9 | 1.7 | 13:44.52 | firefox |
| 15888 | joe | 20 | 0 | 2965132 | 449280 | 101808 | S | 17.9 | 0.7 | 2:41.56 | Isolated Web Co |
| 1868 | joe | 20 | 0 | 4858100 | 509088 | 130992 | S | 8.3 | 0.8 | 1:59.69 | gnome-shell |
| 6124 | joe | 20 | 0 | 3490668 | 799024 | 175960 | R | 7.0 | 1.2 | 3:00.17 | Isolated Web Co |
| 1729 | joe | 20 | 0 | 24.2g | 69640 | 41720 | S | 5.0 | 0.1 | 1:40.57 | Xorg |
| 17920 | joe | 20 | 0 | 741272 | 52504 | 40152 | S | 3.3 | 0.1 | 0:02.08 | gnome-terminal- |
| 938 | root | -51 | 0 | 0 | 0 | 0 | S | 2.0 | 0.0 | 1:12.39 | irq/110-nvidia |
| 7354 | joe | 20 | 0 | 2820044 | 181140 | 105604 | S | 2.0 | 0.3 | 0:29.50 | Isolated Web Co |
| 17682 | joe | 20 | 0 | 2486004 | 132880 | 85512 | S | 2.0 | 0.2 | 0:01.73 | Isolated Web Co |
| 20488 | joe | 20 | 0 | 2413852 | 61436 | 49652 | S | 2.0 | 0.1 | 0:00.06 | Web Content |
| 3300 | joe | 20 | 0 | 2589888 | 146564 | 75932 | S | 1.3 | 0.2 | 0:23.37 | WebExtensions |
| 6525 | joe | 20 | 0 | 85.0g | 342680 | 145632 | S | 1.3 | 0.5 | 2:23.87 | Discord |
| 4291 | joe | 20 | 0 | 382820 | 58228 | 44944 | S | 1.0 | 0.1 | 0:19.24 | RDD Process |
| 5829 | joe | 20 | 0 | 32.6g | 221588 | 84904 | S | 0.7 | 0.3 | 0:55.97 | slack |
| 14 | root | 20 | 0 | 0 | 0 | 0 | I | 0.3 | 0.0 | 0:03.26 | rcu sched |
| 854 | root | 20 | 0 | 2812 | 1116 | 1028 | S | 0.3 | 0.0 | 0:13.32 | acpid |
| 1649 | joe | 9 | -11 | 2882424 | 29512 | 21760 | S | 0.3 | 0.0 | 0:07.03 | pulseaudio |
| 2022 | joe | 20 | 0 | 726844 | 66292 | 50696 | S | 0.3 | 0.1 | 0:00.19 | evolution-alarm |

Process Attributes

- We've seen that a process has an address space
- **Process descriptor** additionally associated with process
- Descriptor exists in kernel space so processes can be managed effectively



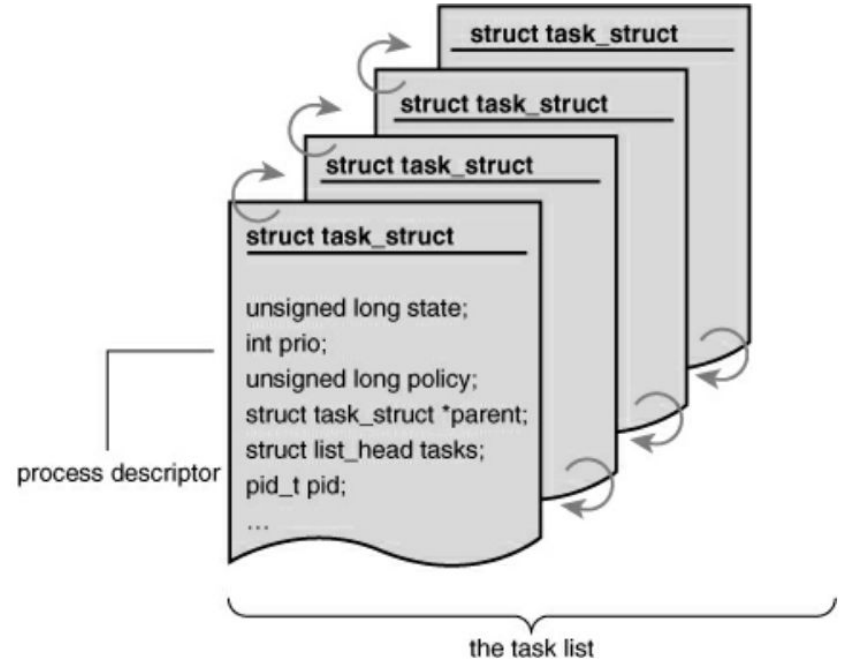
Linux Source Party

[Link to festivities](#)




Process Descriptor

- Instance of `struct task_struct`
 - `mm_struct`
 - `Thread_struct` ([x86](#))
- Motivates viewing the source code of Linux
- Source located at www.kernel.org
- Can find `struct task_struct` defined at line 661 in `include/linux/sched.h` ([link](#))



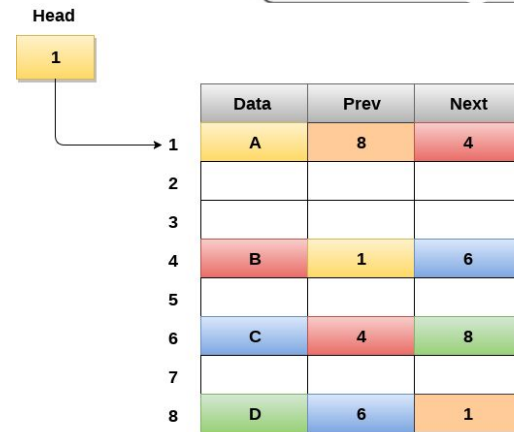
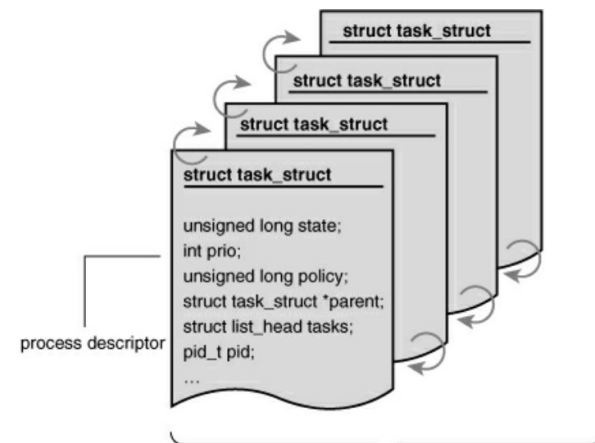
Task List

- Kernel tracks processes in this list
- Circular doubly linked list



Circular Doubly Linked List

<https://www.javatpoint.com/circular-doubly-linked-list>



Memory Representation of a Circular Doubly linked list

<https://www.javatpoint.com/circular-doubly-linked-list>

Implementation of Circular Doubly Linked List in C

```
#include<stdio.h>

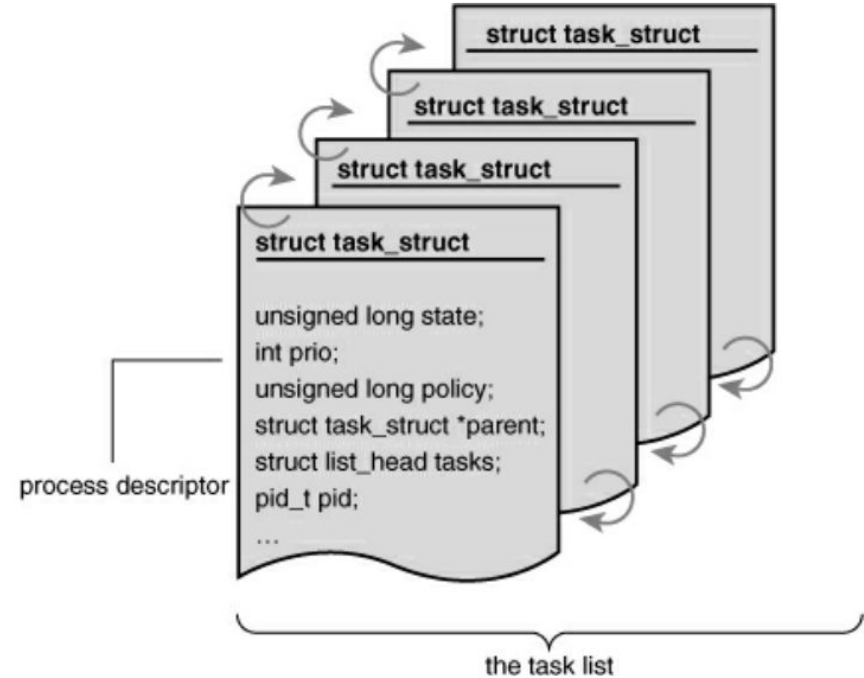
#include<stdlib.h>

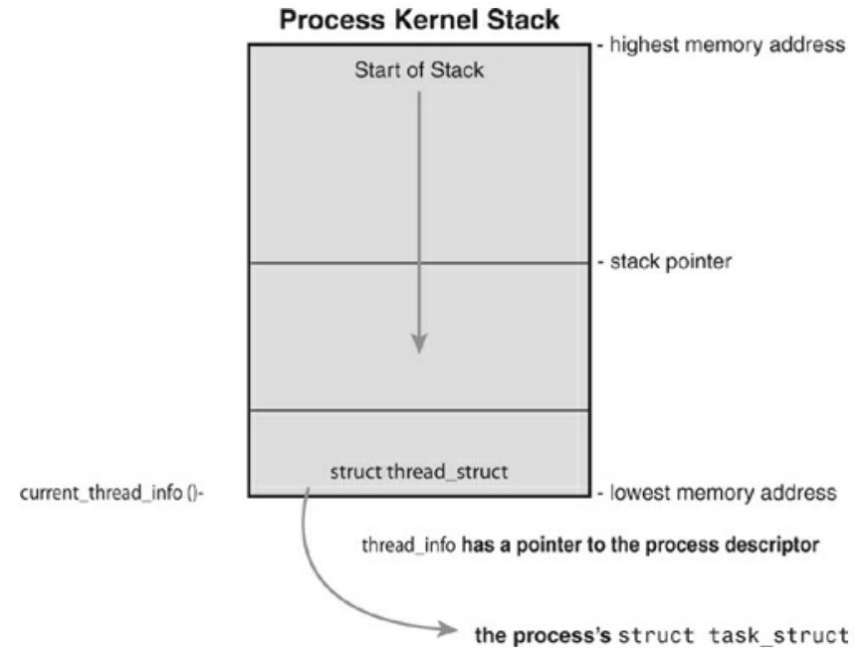
struct node
{
    struct node *prev;
    struct node *next;
    int data;
};

struct node *head;
```

What's Inside `task_struct`? [Process Descriptor]

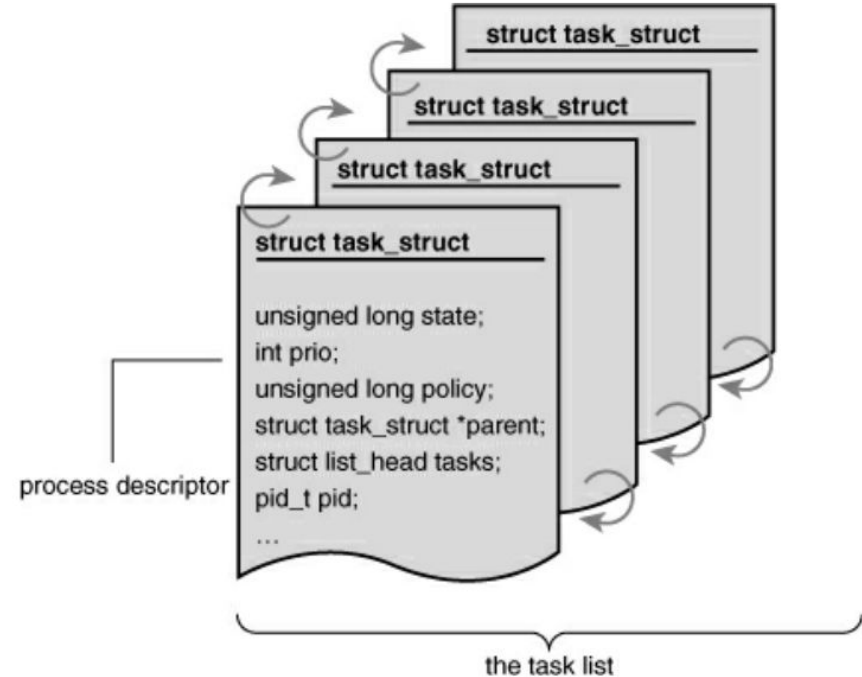
- Open files
- Address space
- Pending signals
- State
- ...



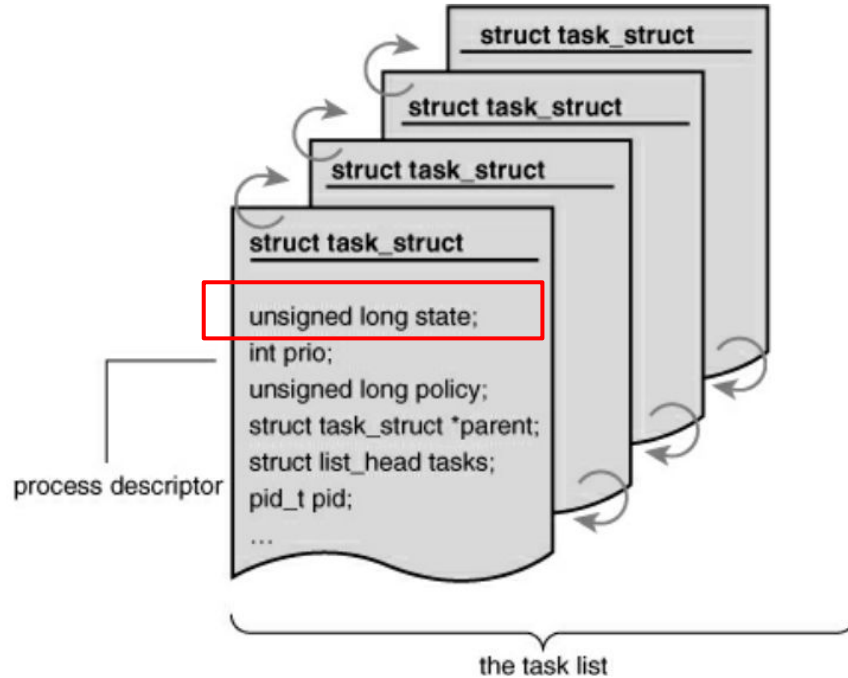


Working With Process Descriptors

- Process identification value (PID) gets assigned to processes
 - In `task_struct`:
 - `pid_t pid;`
 - `pid_t` is an opaque type, usually `int`
- Kernel generally addresses processes' `task_struct` (process descriptor)
- We have a macro ([current](#)) that gets us the presently running process' process descriptor

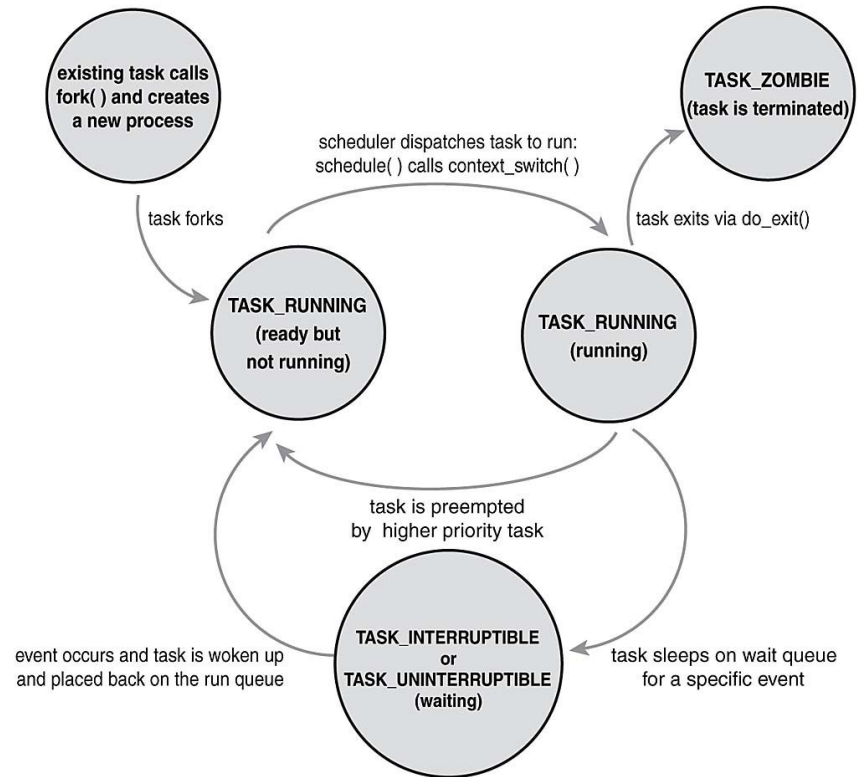


Moving to Process State...



States of Processes ([source](#))

- State stored as a “flag”
- TASK_RUNNING
 - Can be executed
- TASK_INTERRUPTIBLE
 - Sleeping/blocked
 - Awaiting condition or signal
- TASK_UNINTERRUPTIBLE
 - Same as INTERRUPTIBLE, but ignores signals
- Which one could clearly (and does) go wrong?



Altering States of Processes

- Kernel can modify state of process
- [include/linux/sched.h](#)

Linux 4.11 compat: set_task_state() removed

Replace uses of `set_task_state(current, STATE)` with `set_current_state(STATE)`.

In Linux 4.11, torvalds/linux@642fa44, `set_task_state()` is removed.

All spl uses are of the form `set_task_state(current, STATE)`. `set_current_state(STATE)` is equivalent and has been available since Linux 2.2.26.

Furthermore, `set_current_state(STATE)` is already used in about 15 locations within spl. This change should have no impact other than removing an unnecessary dependency.

```
/ include / linux / sched.h
```

```
191  *
192  * Also see the comments of try_to_wake_up().
193  */
194  #define __set_current_state(state_value)
195         WRITE_ONCE(current->__state, (state_value))
196
197  #define set_current_state(state_value)
198         smp_store_mb(current->__state, (state_value))
199
200  /*
201  * set_special_state() should be used for those states
```

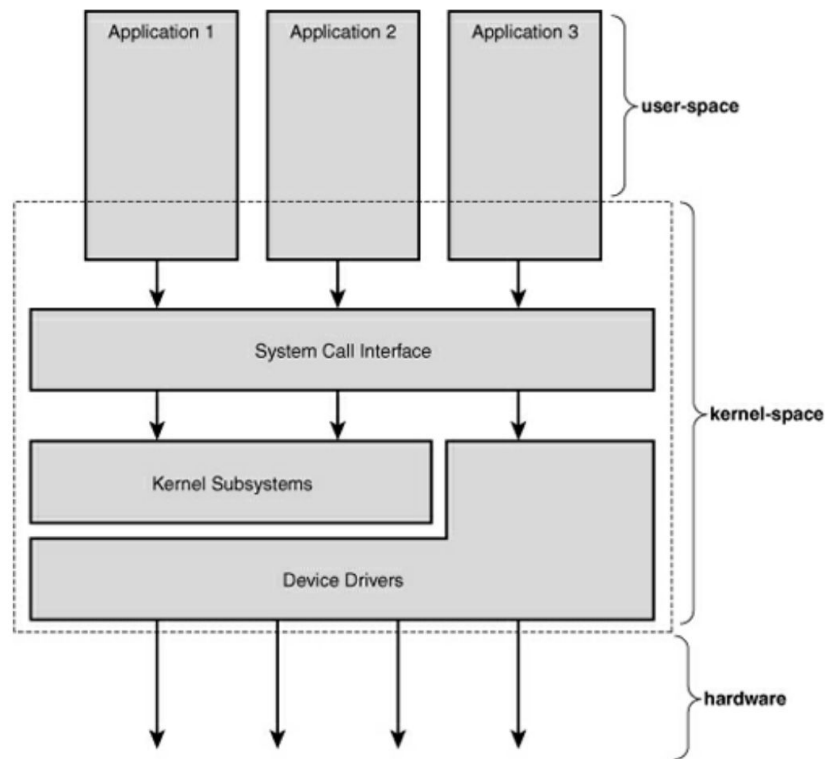
Linux 5.14.1 Source via Bootlin Elixir Cross Referencer

<https://elixir.bootlin.com/linux/v5.14.1/source/include/linux/sched.h#L116>

Process Context

STOPPED HERE

- **User space** is where typical program running happens
- System calls made by the program into **kernel space**
 - Exception could also happen
- ‘Kernel is said to be “executing on behalf of the process” and is in *process context*’
 - Interrupt context is another context
- Legitimate `current` macro
- **System calls** & exception handlers are the exclusive ways into the kernel



Process Hierarchy

- **init**, the ultimate ancestor of all processes
 - Can you guess the PID that it gets?
 - PID: 1
 - **systemd** has replaced this
- 1 parent per process
- 0+ children per process
 - Children referred to as siblings

Process Descriptor (task_struct)

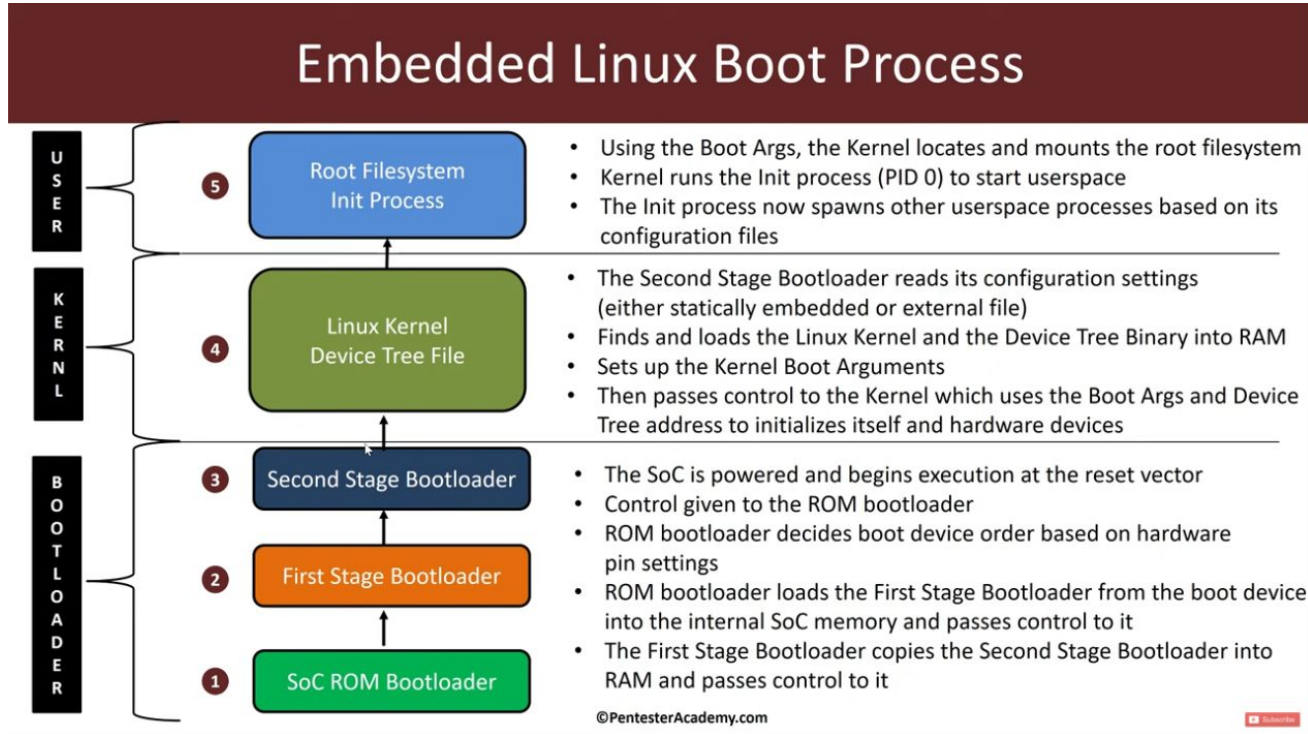
```
/ include / linux / sched.h
878      /*
879      * Pointers to the (original) parent process, youngest child, younger sibling,
880      * older sibling, respectively. (p->father can be replaced with
881      * p->real_parent->pid)
882      */
883
884      /* Real parent process: */
885      struct task_struct __rcu    *real_parent;
886
887      /* Recipient of SIGCHLD, wait4() reports: */
888      struct task_struct __rcu    *parent;
889
890      /*
891      * Children/sibling form the list of natural children:
892      */
893      struct list_head            children;
894      struct list_head            sibling;
895      struct task_struct          *group_leader;
896
```

Linux 5.14.1 Source via Bootlin Elixir Cross Referencer

<https://elixir.bootlin.com/linux/v5.14.1/source/include/linux/sched.h#L878>

```
struct task_struct *my_parent = current->parent;
```

Boot Process (Slide From Pentester Academy TV)



Slide from: Pentester Academy TV, “Embedded Linux Booting Process (Multi-Stage Bootloaders, Kernel, Filesystem)”. (2018) https://www.youtube.com/watch?v=DV5S_ZSdK0s

Generating Processes

- Our first system call: `fork()`
- Duplicates living process
- **Parent** process calls `fork()` to create **child** process
- Notably, `fork()` takes advantage of **copy-on-write**

```
// SPDX-License-Identifier: GPL-2.0-only
```

```
/*
```

```
* linux/kernel/fork.c
```

```
*
```

```
* Copyright (C) 1991, 1992 Linus Torvalds
```

```
*/
```

```
/*
```

```
* 'fork.c' contains the help-routines for the 'fork' system call
```

```
* (see also entry.S and others).
```

```
* Fork is rather simple, once you get the hang of it, but the memory
```

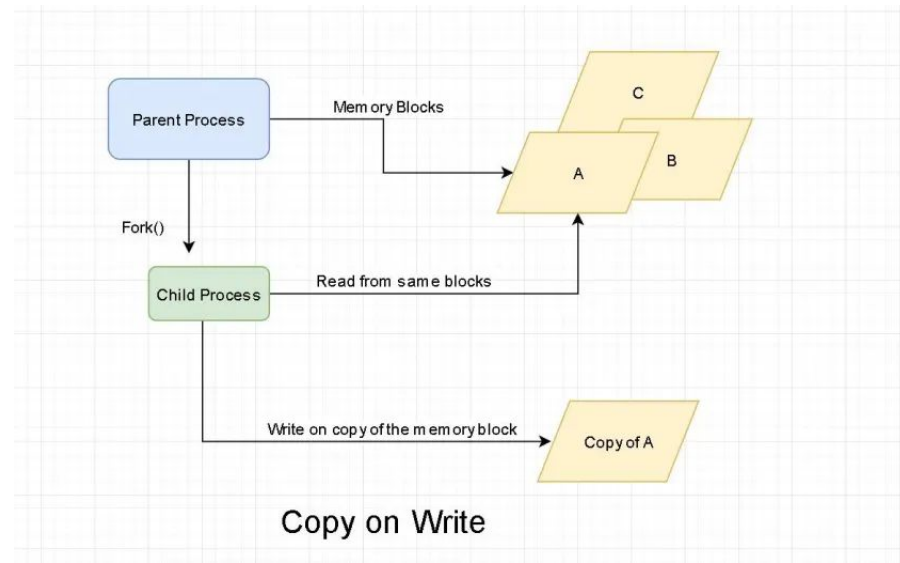
```
* management can be a XXXXXXXXXX. See 'mm/memory.c': 'copy_page_range()'
```

```
*/
```


<https://git.kernel.org/pub/scm/linux/kernel/git/stable/linux.git/tree/kernel/fork.c?h=v5.14.1>

Copy-on-Write (COW)


- Upon generation of child process by parent process, memory pages common to both
- If one process edits a page, it becomes a modified copy of the original page
- Also called **implicit sharing**



Gaurav Yadav, "What is Copy on Write and where is it used?".
<https://www.learnsteps.com/what-is-copy-on-write-and-where-is-it-used/>




Before process P modifies Page 3



After process P modifies Page 3

Copy-on-Write (COW): Reasoning

- Consider if copying happened first
 - Making copies when could just share
 - Copies may also be immediately disregarded
- COW increases efficiency
 - Copy later
 - May never need to copy
- `fork()` is fast
 - Copy page tables & generate PID



Author: Dysprosia. Copyright © [Dysprosia](https://en.wikipedia.org/wiki/File:Virtual_address_space_and_physical_address_space_relationship.png). BSD license. URL to image and license: https://en.wikipedia.org/wiki/File:Virtual_address_space_and_physical_address_space_relationship.png

Virtual Memory - Page Table

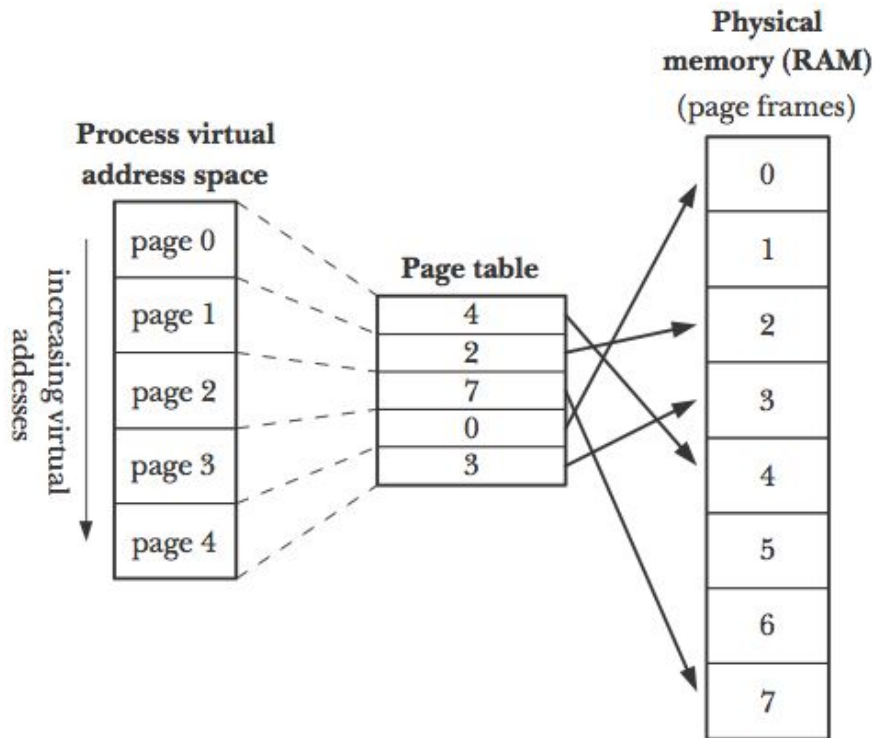
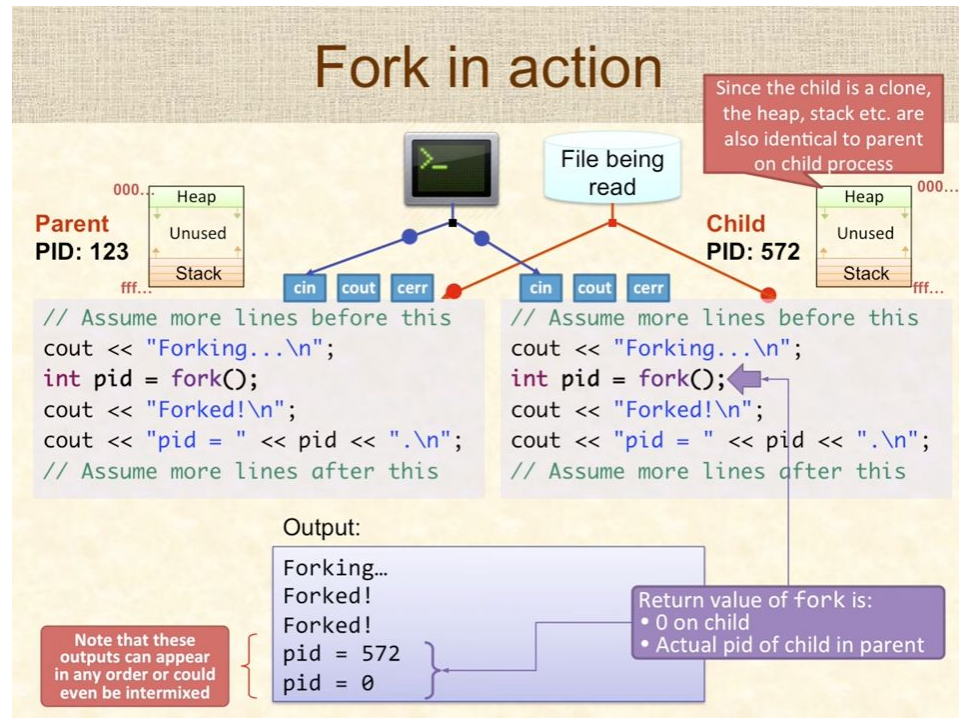


Figure 6-2: Overview of virtual memory

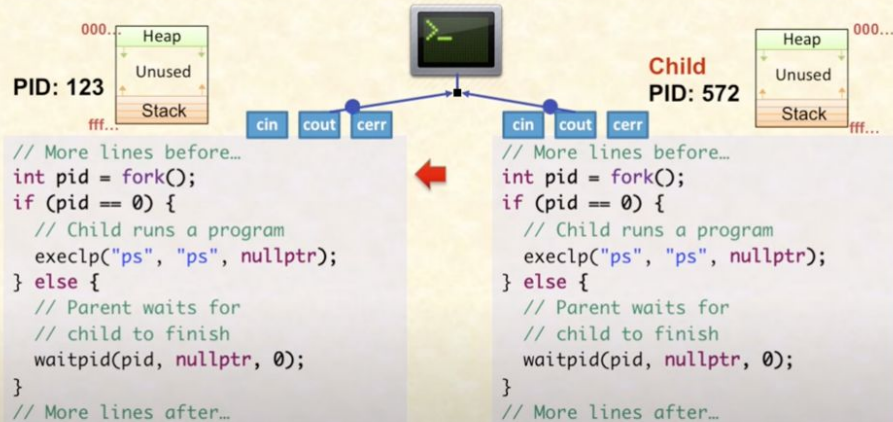
Fork

- /kernel/fork.c
 - Check out the source [here](#)
 - Trace code to `kernel_clone()`
- Returns:
 - 0 if child
 - PID of child if parent

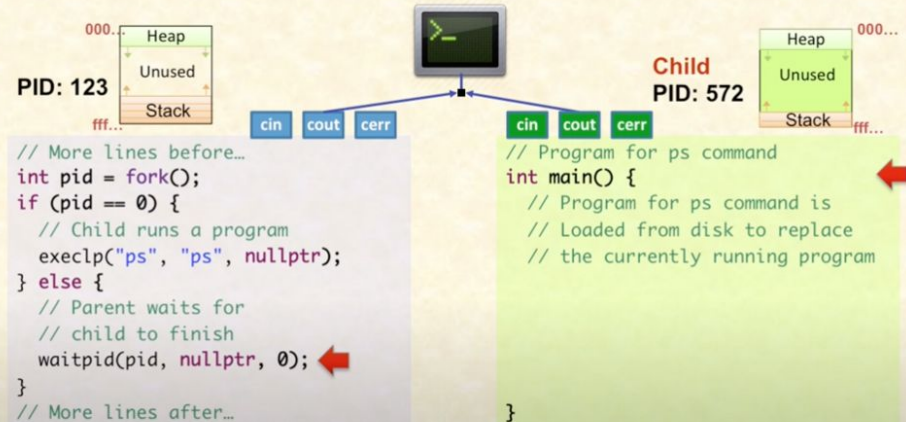


Fork & Exec

Fork + Exec in action



Fork + Exec in action



Kernel Threads

- Kernel-space processes
- Unlike normal processes, **no address space assigned to kernel processes**
 - `mm` literally a NULL pointer
- Schedulable
- Preemptable
- `ps -ef`

The Walking Dead As a Model For Linux Process Death



Where Processes Go To Die (Or How to Get Away With [Process] Murder)

- Process terminates
- Process' resources freed via kernel
- Parent process informed via kernel
- Usually process uses system call `exit()` for termination
 - Could additionally be due to inability to disregard or manage signal/exception
- `do_exit()` takes care of most of the termination process
 - [Link](#) to source in `/kernel/exit.c`
- Final state: `EXIT_ZOMBIE`

Kernel Tree Subdirectories: See Description From David Adams

David Adams, “How to view and browse the linux kernel source?”.
https://linuxhint.com/browse_linux_kernel_source/

Next Lecture:

How Does Linux Schedule Processes?