

ELEC 424/553

Mobile & Embedded Systems

Lecture 5 - How Does Linux Schedule Processes?

The Persistence of Memory by Salvador Dalí (1931). Found at URL hosted by Widewalls:
<https://www.widewalls.ch/magazine/salvador-dali-paintings>

Housekeeping

- Office Hours - Tuesdays 3-4pm - Duncan Hall 2098
 - Raise your hand if this time works for you
- Or by appointment

Housekeeping

- Project 1 posted
- Setting up your Raspberry Pi Zero W
- Due Mon 9/18 11:59pm

fork() demo

Exercise 4

Access Code:

```
#include <stdio.h>
```

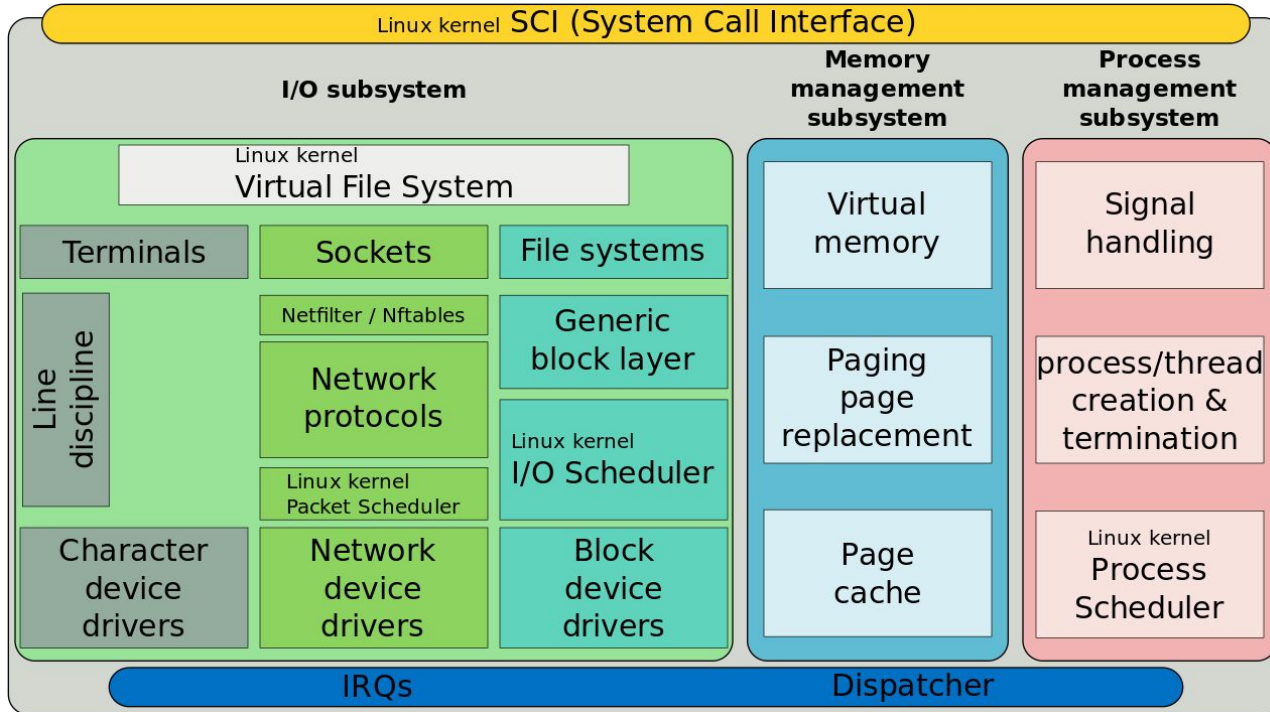
```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
#include <sys/wait.h>
```

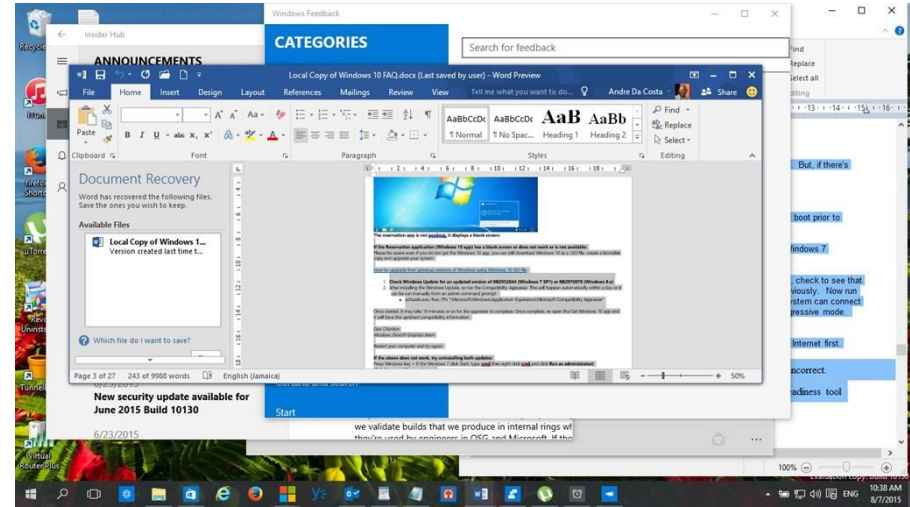
```
Use waitpid(pid, 0, 0);
```

Where Are We In The Linux Kernel?



Scheduling Processes

- Running a bunch of applications
- Multitasking makes it look like OS is simultaneously executing more than 1 application/process
 - Multicore: Simultaneous execution possible
- **Multitasking OS:** Capable of alternating between processes



Andre Da Costa, "How to: manage running programs and virtual desktops using Task View in Windows 10". URL: <https://answers.microsoft.com/en-us/insider/forum/all/how-to-manage-running-programs-and-virtual/17d068b7-5e4a-4351-a019-afa528a81538>

Processes Don't Run All The Time

- **Block/sleep:**
 - Waiting on something to do before running
 - Still taking up memory
 - Not **runnable**
 - Kernel will react when, for example, a key is finally pressed

Two Types of Multitasking

Cooperative



From People.com. Credit: KEVIN MAZUR/WIREIMAGE. URL:
<https://imagesvc.meredithcorp.io/v3/mm/image?url=https%3A%2F%2Fstatic.onecms.io%2Fwp-content%2Fuploads%2Fsites%2F20%2F2019%2F08%2Fswift-2009-5.jpg>

Preemptive



**Linux
Scheduler**

<https://assets.teenvogue.com/photos/570cf9284005974b596c4391/master/pass/Kanye.jpg>

Two Types of Multitasking

Cooperative

Process has the power

Process can give up (yield)
execution whenever it wants

Preemptive

Kernel has the power

Scheduler can come in and
swap out which process is
executing

Timeslice

Non-preemptable time for
an executing process

Original Linux Scheduler

- Linux up to 2.4
 - Higher number of processes/processors would overwhelm scheduler

O(1) Scheduler


- Linux 2.5
- O(1): Independent of # of processes, constant-time scheduling
- Pro: Scaled much better than prior scheduler
- Con: Latency issues
 - Interactive processes
- **Completely Fair Scheduler (CFS)**
introduced in Linux 2.6.23 to take spot of O(1)

Policy


- What is the decision process for when to execute tasks?
- Processes can be:
 - I/O-bound
 - Processor-bound
- Processes have a **priority**
- Timeslices

Process: I/O-Bound or Processor-Bound?

- How is the process investing its time?
 - Running code?
 - Handling I/O requests?
- Processor (CPU)-bound example
 - Infinite loop
 - MATLAB
- I/O examples
 - Solid state drive
 - Keyboard + mouse
 - Network



(a) CPU-bound application



(b) IO-bound application

A Fundamental Tradeoff

Low latency

Responsiveness

vs.

High throughput

vs.

Utilization



From VentureBeat. Image Credit: 20th Century Fox. URL:
<https://venturebeat.com/2016/11/27/minority-report-science-advisor-builds-the-most-awesome-conference-room/>



From: Tom's Hardware. Image credit: Shutterstock. URL:
<https://www.tomshardware.com/news/china-clampdown-crypto-mining-gpu-prices-plunging>

Process Priority

- Consider **priority-based** scheduling
- Ranking
- Round-robin for equal priority
- Priorities in Linux:
 - **nice** (view using `ps -el`)
 - (higher priority) [-20, -19, ..., 18, 19] (lower priority)
 - **Real-time (RT) priority** (view using `ps -eo state,uid,pid,ppid,rtprio,time,comm`)
 - (lower priority) [0, 1, ..., 98, 99] (higher priority)
 - RT processes > normal processes in terms of priority

Timeslice

- Task can execute for a timeslice
- How long is a good standard timeslice?
- A number of OS's shy away from long timeslices
 - E.g. 10 ms
- Processes get dynamic timeslices from CFS
 - Priority will increase or decrease timeslice (relatively)

Scheduling Case Study

- Text editor (I/O-bound)
- Video encoder (processor-bound)
- Text editor barely needs CPU time
 - But you expect it to respond quickly
- Video encoder needs a ton of CPU time
 - But a difference of a second would be unnoticed


Scheduling Case Study (Continued)

- Encoder will get full CPU time when no key pressed
- Key pressed! What happens?
- Editor awakens (blocked->runnable)
- CFS sees editor has used barely any time
- CFS preempts encoder
- CFS runs editor
- Cycle repeats until the end of history

Scheduler Classes

- Modularity of Linux scheduler
- Can use particular scheduler for particular process types
- Priority ranking for process types
- Each class gets a sub run queue (rq) which contains runnable tasks
- Relevant files in `/kernel/sched/`

- [`core.c`](#)
- [`fair.c`](#)



Viresh Kumar, "Fixing SCHED_IDLE". URL:
<https://lwn.net/Articles/805317/>

Scheduler Classes

```
/ kernel / sched / core.c All
5902  * WARNING: must be called with preemption disabled!
5903  */
5904  static void __sched notrace __schedule(bool preempt)
5905  {
5906      struct task_struct *prev, *next;
5907      unsigned long *switch_count;
5908      unsigned long prev_state;
5909      struct rq_flags rf;
5910      struct rq *rq;
5911      int cpu;
5912
5913      cpu = smp_processor_id();
5914      rq = cpu_rq(cpu);
5915      prev = rq->curr;
5916
```

```
/ kernel / sched / sched.h All sym Search
904
905  /*
906   * This is the main, per-CPU runqueue data structure.
907   *
908   * Locking rule: those places that want to lock multiple runqueues
909   * (such as the load balancing or the thread migration code), lock
910   * acquire operations must be ordered by ascending &runqueue.
911   */
912  struct rq {
913      /* runqueue lock. */
914      raw_spinlock_t _lock;
915
916      /*
917       * nr_running and cpu_load should be in the same cacheline because
918       * remote CPUs use both these fields when doing load calculation.
919       */
920      unsigned int uclamp_flags;
921
922      #define UCLAMP_FLAG_IDLE 0x01
923      #endif
924
925      struct cfs_rq cfs;
926      struct rt_rq rt;
927      struct dl_rq dl;
928
```

Linux 5.14.2 Source via Bootlin Elixir Cross Referencer

Left: <https://elixir.bootlin.com/linux/v5.14.2/source/kernel/sched/core.c#L5904>

Right: <https://elixir.bootlin.com/linux/v5.14.2/source/kernel/sched/sched.h#L912>

Next Lecture:

Scheduling Continued