

ELEC 424/553

Mobile & Embedded Systems

Lecture 7 - System Calls



<http://www.cnn.com/2010/TECH/mobile/07/09/cooper.cell.phone.inventor/index.html>

Google Slides Link:

https://docs.google.com/presentation/d/10X89S5DPI8txxseBbGyLh-HgPZX7X8yM_gype2_STa0/edit?usp=sharing

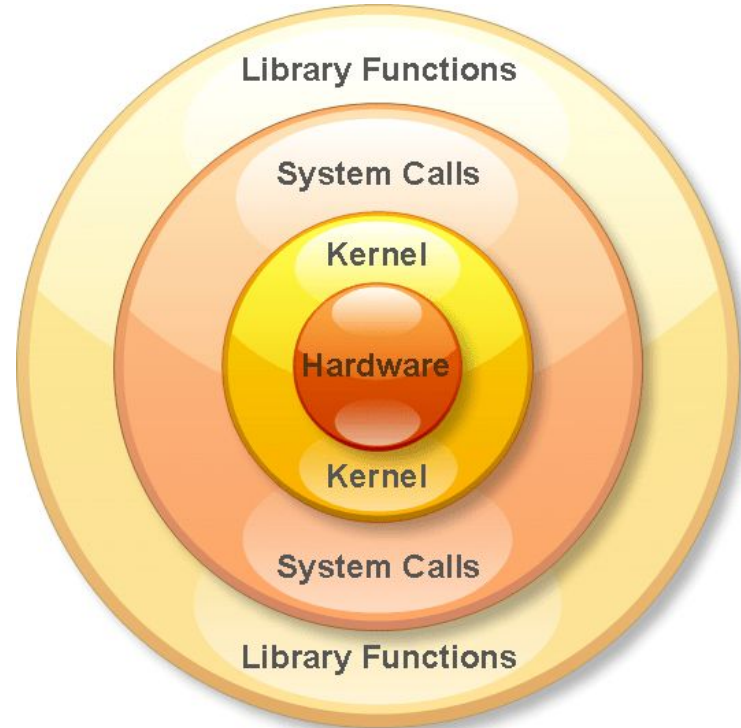
Housekeeping

ARM has hardware support for OS privilege levels similar to x86

Note from prior exercise on factorial code:

- In C, ++n is pre-increment and n++ is post-increment
- For the condition of a while loop, ++n will happen before the condition is checked and n++ will happen after the condition is checked

Soldering session, extra office hours



Balakumaran Kannan, "Anatomy of Linux system call in ARM64". URL: <https://eastrivervillage.com/Anatomy-of-Linux-system-call-in-ARM64/>

Questions In The Past

- **Circular Queue vs. $O(1)$ scheduler**

- Scheduling classes

- **Scheduler overhead**

- $O(\log n)$
- Maybe a few μs for a context switch
- Sources:
 - <https://unix.stackexchange.com/questions/553120/linux-scheduling-overhead-in-orders-of-magnitude>
 - <https://eli.thegreenplace.net/2018/measuring-context-switching-and-memory-overheads-for-linux-threads/>
 - https://events.static.linuxfound.org/sites/events/files/slides/lemoal-nvme-polling-vault-2017-final_0.pdf

- **Processes going from to sleep to awake**

- One response indicates min_vruntime normalization does occur
 - <https://stackoverflow.com/questions/11297285/how-can-the-linux-cfs-scheduler-prevent-a-task-with-a-very-small-vruntime-from-s>

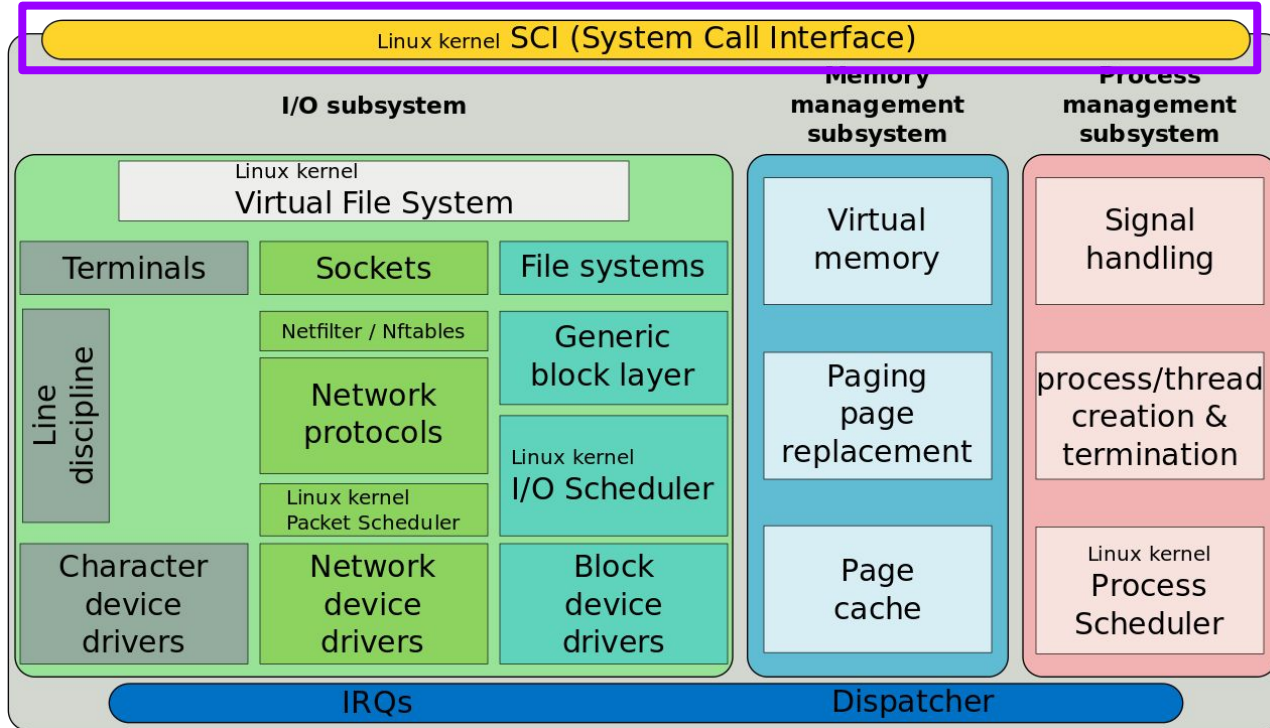
Goals of Today's Lecture

- Enable you to become a **kernel hacker**
- Learn the basics of **system calls**
 - How they're used
 - How they work
 - How to make your own



"Anonymous - Cool Hacker Guy". URL:
<https://i.kym-cdn.com/photos/images/original/001/109/423/cb4.jpg>

Where Are We In The Linux Kernel?

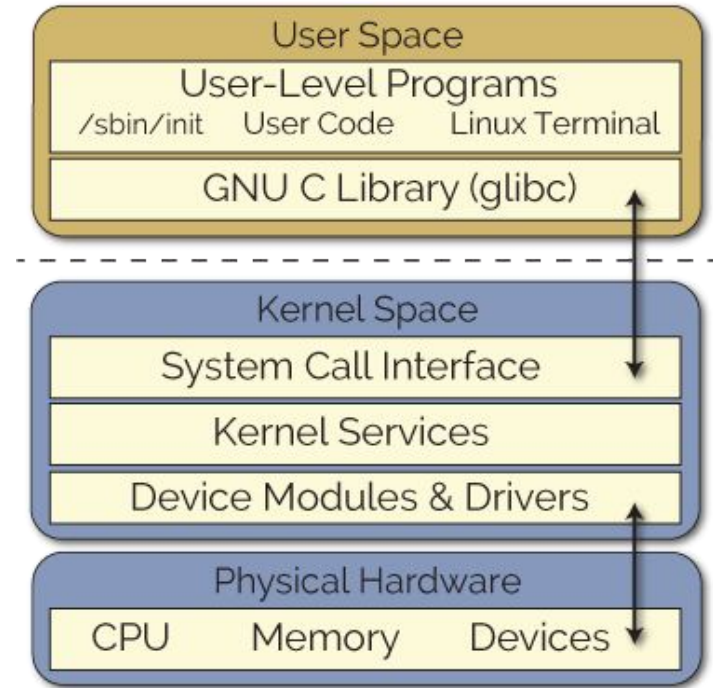


Kernel Space Must Be Protected At All Costs

- User applications have **no direct access** to kernel
- Indirect access via **system calls**
- System calls provide applications with access to **kernel services**
- Kernel acts **on behalf of application**



Rachelle Hampton, "There's Nothing Wrong With Server-Speak. In Fact, It's Polite!". Image source: SeventyFour / Thinkstock.
URL:
<https://slate.com/human-interest/2018/02/food-service-grammar-is-perfectly-polite.html>



Derek Molloy, "Writing a Linux Kernel Module — Part 1: Introduction". URL:
<http://derekmolloy.ie/writing-a-linux-kernel-module-part-1-introduction/#prettyPhoto>

pchdtvr - A Case Study of System Calls In a Program

[Home](#) / [Personal Media Networks](#) [frequent.com](#)

pchdtvr -- record digital television on your PC

pchdtvr is a great program for recording digital television on your PC. It lets you:

1. Set up schedules
2. Record programs instantly
3. Check station signal strength
4. And much more!

<http://frequent.com/pmn/pchdtvr.html>

pchdtvr - A Case Study in Terrible Naming

pchdtvr - a good reason why techs should *never* be allowed to name their products :-)
- user10776 Dec 18 '16 at 8:04

“Can GPL licensed code be close sourced later by the author? [duplicate]”. URL:
<https://softwareengineering.stackexchange.com/questions/98774/can-gpl-licensed-code-be-close-sourced-later-by-the-author/98778>

pchdtvr - A Case Study of License Legal Issues Too

[Home](#) / [Personal Media Networks](#) [frequent.com](#)

pchdtvr -- record digital television on your PC

pchdtvr is a great program for recording digital television on your PC. It lets you:

1. Set up schedules
2. Record programs instantly
3. Check station signal strength
4. And much more!

Best of all, since pchdtvr is licensed under the GNU GPL, nobody can ever take away your right to use it, change it, or give it away!

Download

There are several versions available:

1. [pchdtvr-1.0-frequent](#), the latest and greatest version available.
 1. New % file suffix to attach date to recordings
 2. Append instead of rename for restarted recording
2. [pchdtvr-1.0](#)
3. [pchdtvr-0.9s](#)

Last modified on 26 Jan 2008 by [AO](#)

Copyright © 2020 [Andrew Oliver](#)

<http://frequent.com/pmn/pchdtvr.html>

I have revoked the licensing under the GNU General Public License (herein after referred to as "the GPL") for the atscap version 1.1 codebase, all prior versions of the atscap codebase and all the various release candidates of the atscap codebase.

[...]

As sole author of both the atscap and the pchdtvr codebases, the licenses under the GPL were granted at my sole discretion and the licenses under the GPL are now hereby revoked at my sole discretion.

[...]

If you are currently using the atscap or pchdtvr packages, or any part thereof, it is in your best interest to remove the software from your system(s) and destroy all copies in your possession.

[...]

If you have incorporated the atscap or pchdtvr codebase, or any part thereof, into any of your projects, it is in your best interest to remove any and all of my code from your project(s).

<https://blog.jasonantman.com/2008/01/im-sorry-but-i-dont-like-this-contract-anymore/>

pchdtvr - Not a Totally Satisfying Conclusion in Regards to Legal Issues

I've never used this package (apparently its used for HDTV scheduling/recording on Linux), but this link on Slashdot caught my eye: http://sourceforge.net/developer/diary.php?diary_id=26407&diary_user=147583. Apparently, the developer of this software package is seeking to revoke the GPL license not just for his current code, but his past code/package as well. I have a difficult time believing that this is possible, but I'm sure we will soon find out. My guess is that this guy is productizing his software and has a good idea who is currently using, selling and distributing his source so there will likely be some kind of legal challenge to the GPL as well. It's always interesting to see how these kind of things play out in the U.S. courts which can sometimes be a little schizophrenic, though I'd have a difficult time believing that this type of retroactive license change is actually possible.

User: reeset, "atscap and pchdtvr GPL revoked or can it be".
URL: <https://blog.reeset.net/archives/493>

The screenshot shows a Stack Exchange question on the 'Software Engineering' site. The question title is 'Can GPL licensed code be close sourced later by the author? [duplicate]'. The question body contains two paragraphs. The first paragraph discusses the GPL license and the possibility of revoking it. The second paragraph mentions a 'semi-famous case' where a developer tried to revoke the GPL from his code and was 'flamed right off the net'. The question is marked as 'duplicate' and has a score of 6. The answer section shows a response from user 'raptortech97' dated Feb 6 '15 at 14:44, stating that 'Flamed right off the net' doesn't sound like a legal procedure. Another answer from user 'pchdtvr' dated Dec 18 '16 at 8:04 is partially visible.

StackExchange Search on Software Engineering... Log in Sign u

Home PUBLIC Questions Tags Users Unanswered FIND A JOB Jobs Companies TEAMS

Stack Overflow for Teams – Collaborate and share knowledge with a private group.

If you own the code and you have GPL'd it, you can't take it back. The best you can do is grant yourself (or someone else) the right to develop the code further under a closed model.

4 There was a semi-famous case of a few years back, with some software I actually use called [pchdtvr](#). The guy who wrote it was a pretty good guy, but he tried to revoke the GPL from his code and basically got flamed right off the net. (Angry people basically forced him to shut down his domain, and he disappeared.) You could search for the package or search for his email, inkling@nop.org.

So, closing your code and putting it up for sale is no problem. Trying to stop people from giving away your GPL'd code? Big problem.

Share Improve this answer Follow edited Feb 6 '15 at 15:07 user40980 answered Aug 5 '11 at 2:17 Norman Ramsey 2,621 ●1 ●17 ●15

6 "Flamed right off the net" doesn't sound like a legal procedure – raptortech97 Feb 6 '15 at 14:44

[pchdtvr](#) – a good reason why techs should never be allowed to name their products :-)
– user10776 Dec 18 '16 at 8:04

"Can GPL licensed code be close sourced later by the author? [duplicate]". URL: <https://softwareengineering.stackexchange.com/questions/98774/can-gpl-licensed-code-be-close-sourced-later-by-the-author/98778>

pchdtvr - GPL Means We Still Have the Source

```
pchdtvr.c
19
20 /*****
21  *
22  * pchdtvr.c    (c) Copyright 2004-2006 by inkling@nop.org
23  *    ATSC Transport Stream Capture Application Program
24  *    for the pchDTV Inc., HD-2000 and HD-3000 Terrestrial DTV adapters
25  *
26  * Watch your DTV shows WHEN, WHERE and HOW you want to watch them.
27  *
28  * ***** pchdtvr is ONLY for Digital Terrestrial (Over The Air) *****
29  *
30  * pchdtvr is free software; you may only redistribute it and/or modify
31  * it under the terms of the GNU General Public License Version 2, as
32  * published by the Free Software Foundation.
33  *
34  * pchdtvr is distributed to you in the hope that it will be useful,
35  * but WITHOUT ANY WARRANTY OR SUPPORT; without even the implied warranty of
36  * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. Please see the
37  * GNU General Public License for more details.
38  *
39  * You should have received a copy of the GNU General Public License Version 2
40  * along with this program; if not, write me or the Free Software Foundation,
41  * Inc., at 51 Franklin Street, Fifth Floor, Boston, MA 02110-1301 USA
42  *
43  *****/
```

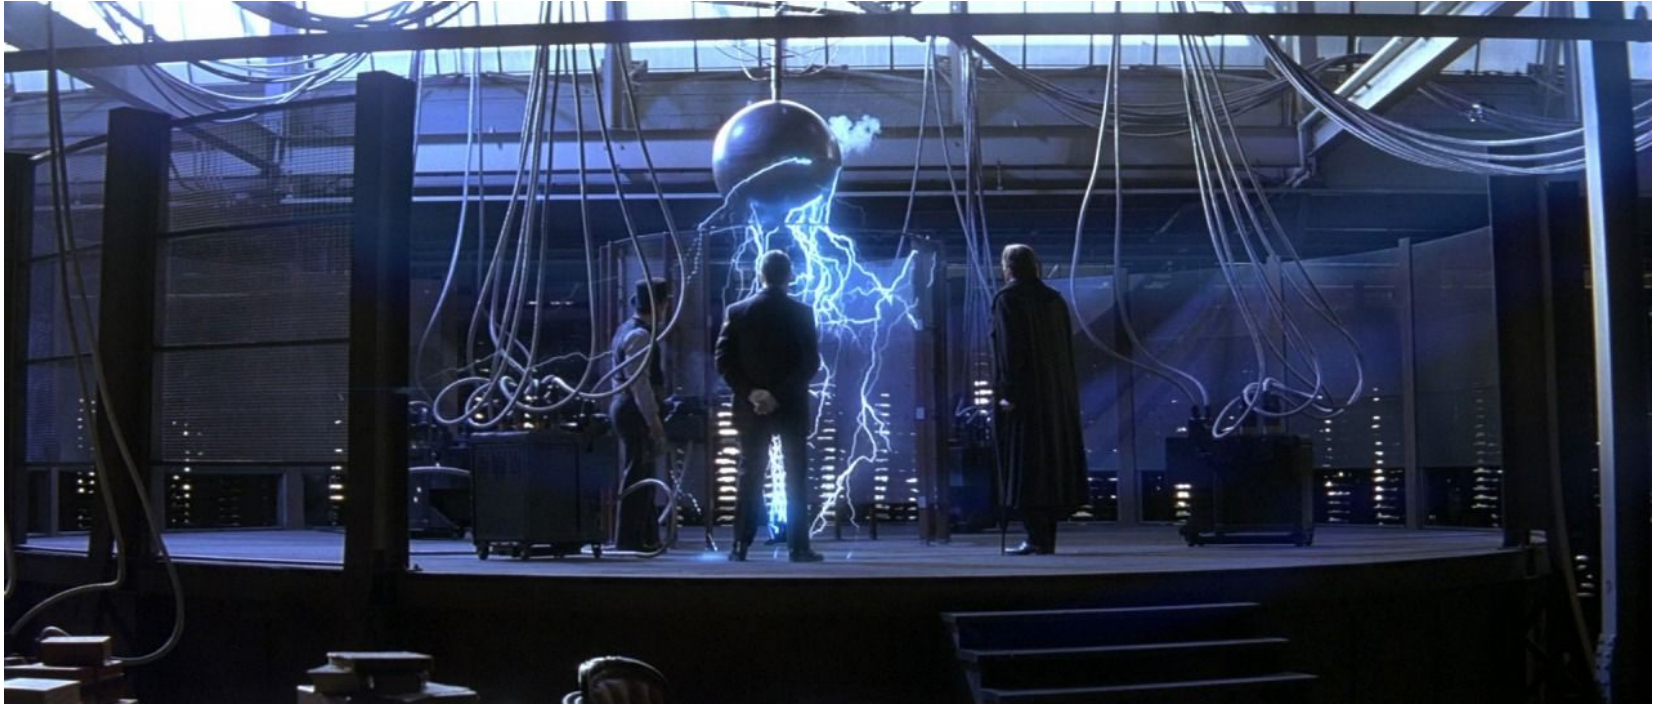
pchdtvr - We Can Find Our Old Friend `fork()`

```
pchdtvr.c
22250     if (0 != arg_detach)
22251     {
22252     pid_t fp;
22253     fp = fork();
22254     if (0 == fp)
22255     {
22256         /* new forked path, will fall thru */
22257         pid_m = getpid();
22258         asyslog( LOG_INFO,
22259             "daemon forked, main() is now pid %d", pid_m );
22260         /* fall thru to below */
22261     } else {
22262         /* old forker must die right here */
22263         console_exit(0);
22264     }
```

Remember: `pid_t` is an opaque type (not known by application)

`fork()` called

Christopher Nolan's Explanation of `fork()` Return Value



pchdtvr - We Can Find Our Old Friend `fork()`

```
pchdtvr.c
22250     if (0 != arg_detach)
22251     {
22252     pid_t fp;
22253     fp = fork();
22254     if (0 == fp)
22255     {
22256         /* new forked path, will fall thru */
22257         pid_m = getpid();
22258         asyslog( LOG_INFO,
22259             "daemon forked, main() is now pid %d", pid_m );
22260         /* fall thru to below */
22261     } else {
22262         /* old forker must die right here */
22263         console_exit(0);
22264     }
```

Remember: `pid_t` is an opaque type

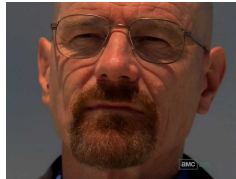
`fork()` called

Am I the **child process**?

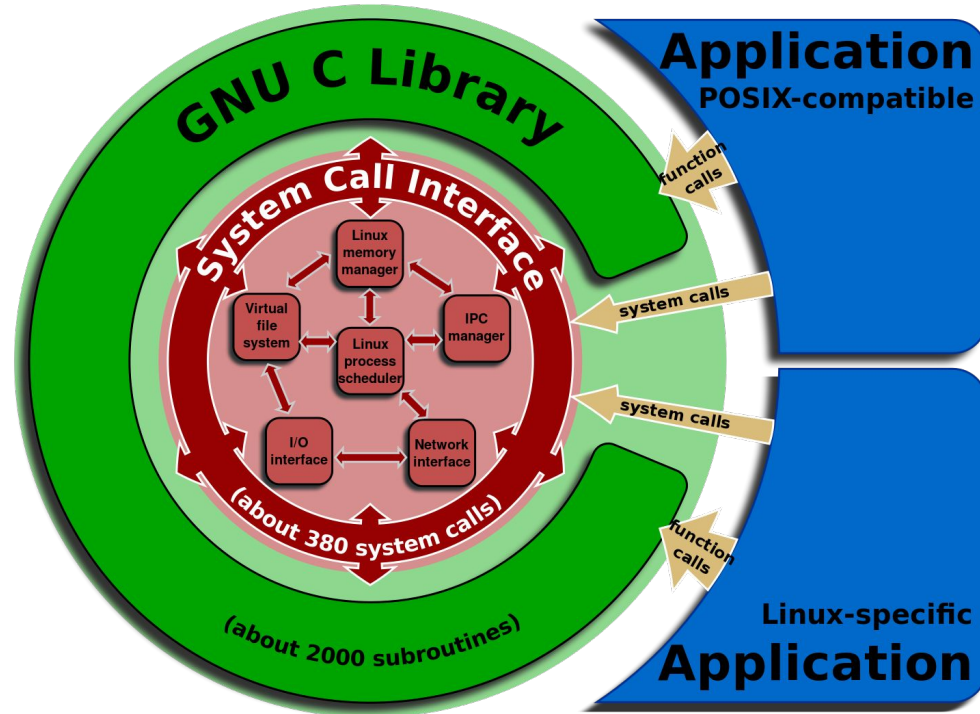
What is my **PID**?

“say my name”

I’m the **parent process**
Death imminent



Using System Calls **Directly** & **Indirectly**



Indirectly Using System Call Via **glibc**

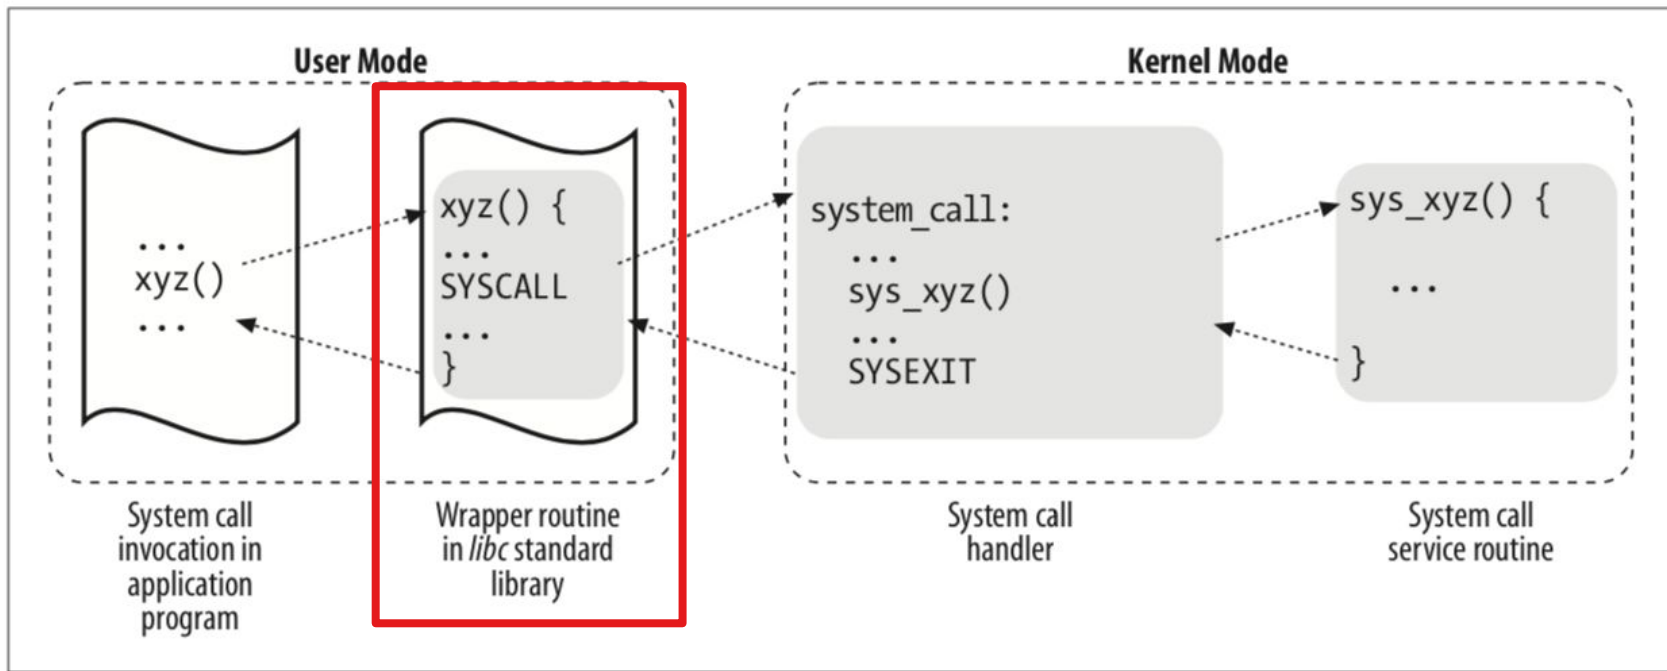
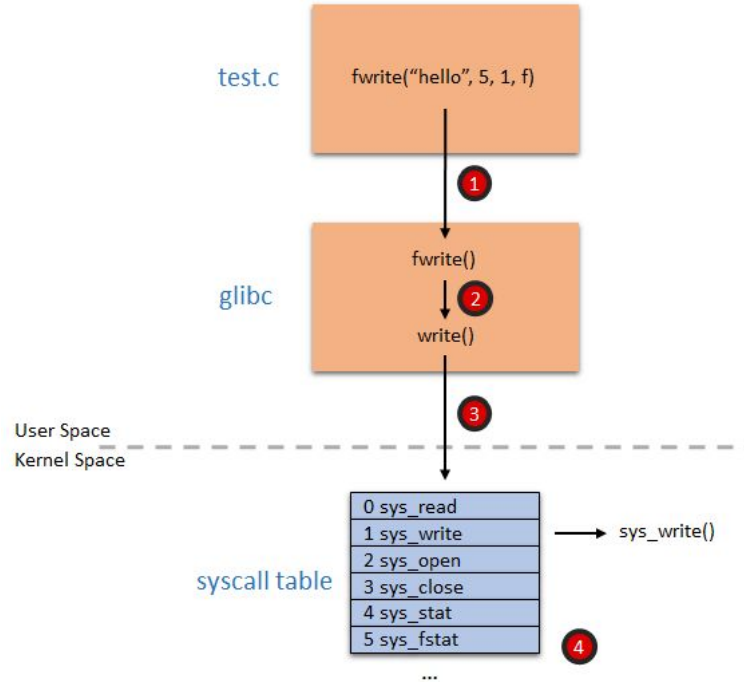


Figure 10-1. Invoking a system call

Example of `write()` System Call Being Used



fork () Does Not Call fork()

fork(2) — Linux manual page

[NAME](#) | [SYNOPSIS](#) | [DESCRIPTION](#) | [RETURN VALUE](#) | [ERRORS](#) | [CONFORMING TO](#) | [NOTES](#) | [EXAMPLES](#) | [SEE ALSO](#) | [COLOPHON](#)

FORK(2)

Linux Programmer's Manual

FORK(2)

NAME [top](#)
fork – create a child process

C library/kernel differences

Since version 2.3.3, rather than invoking the kernel's `fork()` system call, the glibc `fork()` wrapper that is provided as part of the **NPTL threading** implementation invokes `clone(2)` with flags that provide the same effect as the traditional system call. (A call to `fork()` is equivalent to a call to `clone(2)` specifying *flags* as just `SIGCHLD`.) The glibc wrapper invokes any fork handlers that have been established using `pthread_atfork(3)`.

Remember when?...

- 1983: UNIX System V released
 - **Richard Stallman also started GNU Project**
 - Started developing free programs functionally equivalent to UNIX programs
- The UNIX wars (late 80s to mid 90s)
 - What is UNIX? Many versions were floating around
- Efforts for peace were made
 - 1985: AT&T introduced the System V Interface Definition (SVID)
 - **1988: IEEE POSIX (Portable Operating System Interface)**
 - AT&T worked with others to merge other UNIX versions into System V
 - The Open Software Foundation countered with OSF/1
 - All the while, more commercial distributions of UNIX appeared that expanded beyond AT&T's standard



Richard Stallman

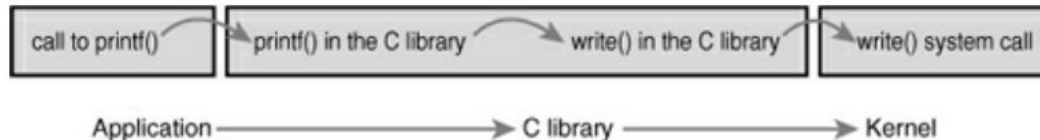
From the cover of: *Free as in Freedom*, Sam Williams.
https://www.amazon.com/Free-Freedom-Richard-Stallmans-Software-dp-0596002874/dp/0596002874/ref=mt_other?_encoding=UTF8&me=&qid=1630343742

POSIX - Portable Operating System Interface

- **POSIX** encourages compatibility among OS's for applications via **set of standards**
- UNIX Wars - Guarantee application worked across UNIX flavors
- Standardizes:
 - Shell
 - **API (application programming interface) to kernel**
- Many things are not standardized
 - GUIs
- This does not talk about **binary compatibility**

POSIX - API Specifics

- Common API enables applications to access kernel services across UNIX and UNIX-like systems
- POSIX talks about **System Interfaces**:
 - “the functions, macros, and external variables to support applications portability at the C-language source level” [IEEE & The Open Group: [link](#)]
- Implementation unspecified - can include:
 - System calls via the C library
 - Staying *entirely* within the C library (user space)
- **glibc** - The GNU C Library
 - C standard library realization by GNU Project



POSIX - System Interface Examples

- List of POSIX system interfaces available on The Open Group's site [here](#)
 - Again, not the same, but we have a list of Linux system calls [here](#)
- System interface: [fork\(\)](#)
 - System call: [fork\(\)](#)
 - But we know glibc will not use the `fork()` system call, it will use the `clone()` system call
 - Note: no `clone()` system interface
- System interface: [fopen\(\)](#)
 - System call: [open\(\)](#)
- System interface: [sqrt\(\)](#)
 - System call: Assumably none

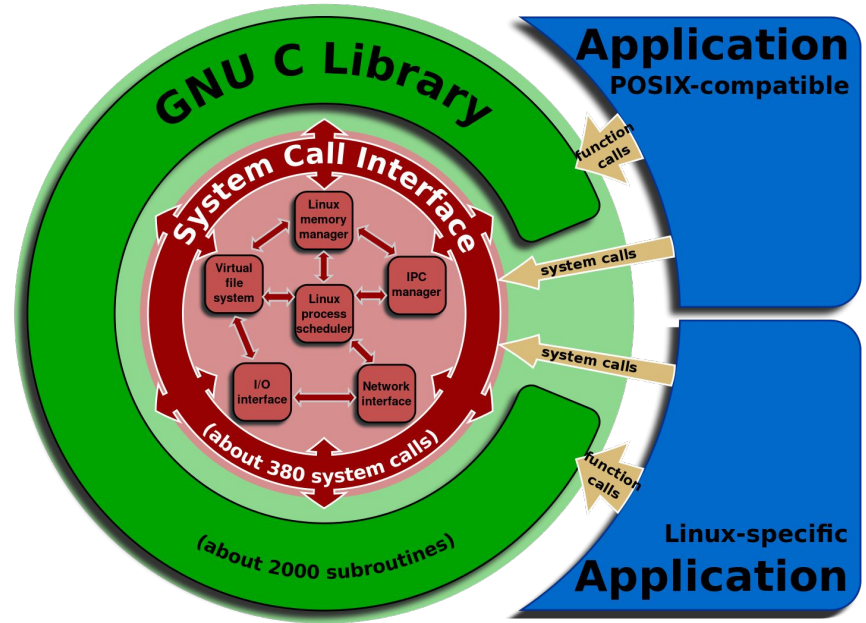
Using System Call **Directly** or **Indirectly**

```
#include <syscall.h>
#include <unistd.h>
#include <stdio.h>
#include <sys/types.h>

int main(void) {

    long ID1, ID2;
    /*-----*/
    /* direct system call */
    /* SYS_getpid (func no. is 20) */
    /*-----*/
    ID1 = syscall(SYS_getpid);
    printf ("syscall(SYS_getpid)=%ld\n", ID1);

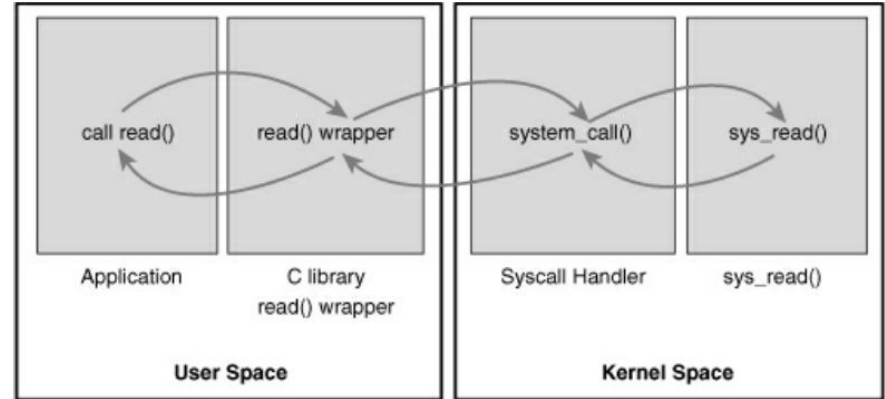
    /*-----*/
    /* "libc" wrapped system call */
    /* SYS_getpid (Func No. is 20) */
    /*-----*/
    ID2 = getpid();
    printf ("getpid()=%ld\n", ID2);
    return(0);
}
```



ScotXW. CC Attribution-Share Alike 3.0 Unported license ([link](https://commons.wikimedia.org/wiki/File:Linux_kernel_System_Call_Interface_and_glibc.svg)). URL: https://commons.wikimedia.org/wiki/File:Linux_kernel_System_Call_Interface_and_glibc.svg

How Does This Really Work? (x86 Edition)

- Cause a **software interrupt**: exception/trap/fault
- **Kernel mode** activated
- Exception handler (**system call handler**) runs
- Which system call?
 - Use **eax register** to give system call # to kernel
 - **eax** read by handler



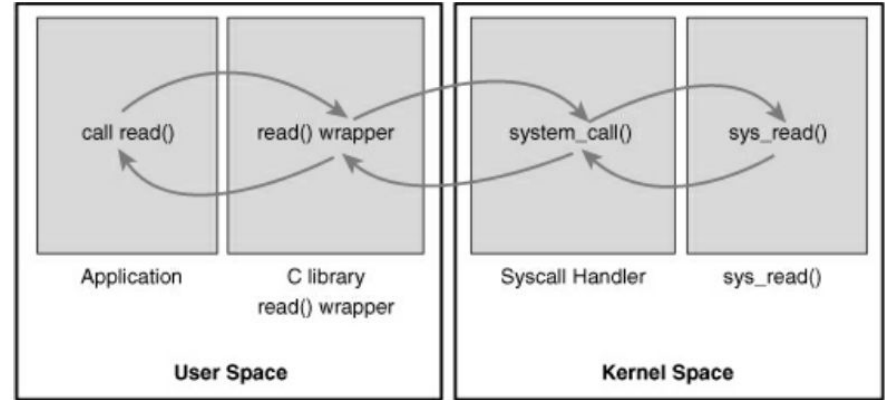
eax

System Call Numbers (x86-64)

/ arch / x86 / entry / syscalls / syscall_64.tbl				All sy ▾	Search Identifier 🔍
1	#				
2	#	64-bit system call numbers and entry vectors			
3	#				
4	#	The format is:			
5	#	<number> <abi> <name> <entry point>			
6	#				
7	#	The __x64_sys_*() stubs are created on-the-fly for sys_*() system calls			
8	#				
9	#	The abi is "common", "64" or "x32" for this file.			
10	#				
11	0	common	read	sys_read	
12	1	common	write	sys_write	
13	2	common	open	sys_open	
14	3	common	close	sys_close	
15	4	common	stat	sys_newstat	
16	5	common	fstat	sys_newfstat	
17	6	common	lstat	sys_newlstat	
18	7	common	poll	sys_poll	

Parameter Passing & Return Value

- Use mechanism similar to syscall # transmission
- Place parameters in **registers**
- Shown registers are for 5 arguments
- If more needed, put user-space pointer in 1 register
 - No other registers used
- How do we get the return value?
 - **eax**




4 Simple Steps to Making Your Very Own System Call

1. What is the point of its existence?
2. Write the code & add to file
3. Add call to the system call table
4. Compile

1. What Is The Point of The System Call?

- Linux philosophy: syscall does one thing
- Determine:
 - Input arguments
 - Return value
 - Error codes
- Unix motto: “Provide mechanism, not policy”
- Portability, robustness
 - Fundamental syscalls have been helpful for 30 years!

2. Write The System Call Code: Parameter Verification

- We're in **kernel space**, tread carefully
- **File I/O**: Is file descriptor legitimate?
- **Processes**: Is PID legitimate?
- **Permissions**: Is the process allowed to see this?
- **Pointers** 
 - Is this really pointing to user space?
 - Is this really pointing to process's space?
 - Is the memory readable/writable/executable?

2. Example of System Call Code

```
SYSCALL_DEFINE1(stephen, char *, msg)
{
    char buf[256];
    long copied = strncpy_from_user(buf, msg, sizeof(buf));
    if (copied < 0 || copied == sizeof(buf))
        return -EFAULT;
    printk(KERN_INFO "stephen syscall called with \"%s\\n\"", buf);
    return 0;
}
```

Stephen Brennan, "Tutorial - Write a System Call". URL: <https://brennan.io/2016/11/14/kernel-dev-ep3/>

Include syscall code in kernel file

- E.g. `sys.c`

3. Let's Make It Official: Add to Syscall Table

- Append call to end of syscall table
 - Number defined by ranking in table
 - Architecture-specific

Add to `syscall_64.tbl`

/ arch / x86 / entry / syscalls / <code>syscall_64.tbl</code>				All syml▼	Search
361	437	common	openat2	sys_openat2	
362	438	common	pidfd_getfd	sys_pidfd_getfd	
363	439	common	faccessat2	sys_faccessat2	
364	440	common	process_madvise	sys_process_madvise	
365	441	common	epoll_pwait2	sys_epoll_pwait2	
366	442	common	mount_setattr	sys_mount_setattr	
367	443	common	quotactl_fd	sys_quotactl_fd	
368	444	common	landlock_create_ruleset	sys_landlock_create_ruleset	
369	445	common	landlock_add_rule	sys_landlock_add_rule	
370	446	common	landlock_restrict_self	sys_landlock_restrict_self	
371	447	common	memfd_secret	sys_memfd_secret	
372					

4. Final Step

Compile and cross your fingers

Test Ride Your Beautiful New Syscall

- **glibc** contains the accepted system calls
- Yours is brand new!
- Use `syscall()` function

```
/*
 * Test the stephen syscall (#329)
 */
#define _GNU_SOURCE
#include <unistd.h>
#include <sys/syscall.h>
#include <stdio.h>

/*
 * Put your syscall number here.
 */
#define SYS_stephen 329

int main(int argc, char **argv)
{
    if (argc <= 1) {
        printf("Must provide a string to give to system call.\n");
        return -1;
    }
    char *arg = argv[1];
    printf("Making system call with \"%s\".\n", arg);
    long res = syscall(SYS_stephen, arg);
    printf("System call returned %ld.\n", res);
    return res;
}
```

Stephen Brennan, "Tutorial - Write a System Call". URL:
<https://brennan.io/2016/11/14/kernel-dev-ep3/>