

**ELEC 424/553**

# **Mobile & Embedded Systems**

**Lecture 10 - Writing  
Your Own Char Driver**

Image URL:

<https://thrivingmarriages.com/not-sure-what-to-say-here-are-100-love-letter-prompts/>

# Why Character Devices?



# NVIDIA Open Source Driver -

<https://github.com/NVIDIA/open-gpu-kernel-modules>

☰ README.md

## NVIDIA Linux Open GPU Kernel Module Source [↗](#)

This is the source release of the NVIDIA Linux open GPU kernel modules, version 535.113.01.

### How to Build [↗](#)

To build:

```
make modules -j$(nproc)
```



To install, first uninstall any existing NVIDIA kernel modules. Then, as root:

```
make modules_install -j$(nproc)
```



Note that the kernel modules built here must be used with GSP firmware and user-space NVIDIA GPU driver components from a corresponding 535.113.01 driver release. This can be achieved by installing the NVIDIA GPU driver from the .run file using the `--no-kernel-modules` option. E.g.,

```
sh ./NVIDIA-Linux- [... ].run --no-kernel-modules
```



# Housekeeping

- **Assignment 1 - Kernel Hacking**

- Add your own system call and use it
- Due next Wednesday October 4th at 11:59pm
- NOTE: If you prefer to submit a video demo, please upload it to the Box folder that I shared with your Rice email.
  - **It is your responsibility to let me know 48 hours before the deadline if you do not have access to the folder, otherwise it is not my responsibility if you have to submit late.**
  - Video files submitted must be named as the following (replacing words as appropriate):  
ELEC424\_Assignment1\_Firstname\_Lastname\_netID

- **Again: Bring Raspberry Pi to every future class**

- Let me know after class today if this is an issue

# Don't Do This At Home!

- Recommend not using your primary laptop and operating system
  - However, you are unlikely to encounter issues
  - Just be sure to save anything else running on your computer and be ready for a random crash; Don't hold me liable for this!
- Reportedly does not work on WSL, and will not work on CLEAR (you need sudo)
- Use Raspberry Pi, VirtualBox, or some other way of protecting your main system and files

# The Code You'll See Today Is a Combination of:

- Derek Molloy's (Dr. Derek Molloy, School of Electronic Engineering, Dublin City University, Ireland) excellent work here:  
<http://derekmolloy.ie/writing-a-linux-kernel-module-part-1-introduction/>  
<http://derekmolloy.ie/writing-a-linux-kernel-module-part-2-a-character-device/>
- Corbet, Rubini, & Kroah-Hartman, Linux Device Drivers, 3rd Ed. URL:  
<https://lwn.net/Kernel/LDD3/>
- My own **craziness**

# init & exit

init

- insmod runs this

exit

- rmmod runs this

# Log Levels (KERN\_INFO) For `printk()`

```
#include <linux/module.h>
#include <linux/init.h>
#include <linux/kernel.h>

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Abraham Lincoln");
MODULE_DESCRIPTION("Greatest module in the world!");
MODULE_VERSION("0.000001");

static int __init hello_init(void){
|   printk(KERN_INFO "Oh hi mark\n");
    return 0;
}

static void __exit hello_exit(void){
|   printk(KERN_INFO "sad, but still love you\n");
}

module_init(hello_init);
module_exit(hello_exit);
```



# Adding a Module Parameter

- `insmod` can include multiple parameter values
- Example from LDD book:

```
insmod hello.ko howmany=10 whom="Mom"
```

- In the module, we use a macro: `module_param`
  - See `moduleparam.h` [here](#)
- Three inputs to `module_param`:
  - Variable name
  - Variable type
  - Permissions mask

```
/ include / linux / moduleparam.h All symbols ▼
101 /**
102  * module_param - typesafe helper for a module/cmdline parameter
103  * @name: the variable to alter, and exposed parameter name.
104  * @type: the type of the parameter
105  * @perm: visibility in sysfs.
106  *
107  * @name becomes the module parameter, or (prefixed by KBUILD_MODNAME and a
108  * ".") the kernel cmdline parameter. Note that - is changed to _, so
109  * the user can use "foo-bar=1" even for variable "foo_bar".
110  *
111  * @perm is 0 if the variable is not to appear in sysfs, or 0444
112  * for world-readable, 0644 for root-writable, etc. Note that if it
113  * is writable, you may need to use kernel_param_lock() around
114  * accesses (esp. charp, which can be kfree'd when it changes).
115  *
116  * The @type is simply pasted to refer to a param_ops_##type and a
117  * param_check_##type: for convenience many standard types are provided but
118  * you can create your own by defining those variables.
119  *
120  * Standard types are:
121  *   byte, hexint, short, ushort, int, uint, long, ulong
122  *   charp: a character pointer
123  *   bool: a bool, values 0/1, y/n, Y/N.
124  *   invbool: the above, only sense-reversed (N = true).
125  */
126 #define module_param(name, type, perm) \
127     module_param_named(name, name, type, perm)
```

Linux Source via Bootlin Elixir Cross Referencer

<https://elixir.bootlin.com/linux/latest/source/include/linux/moduleparam.h>

# Permissions Mask

- Our third field for `module_param`
- Check out `stat.h` - [link](#)
  - Also check out other `stat.h` file [here](#) (not pictured)
- Header gives definitions for permissions macros
- `S_IRUGO` - Anyone can read (can't modify)
- `S_IRUGO | S_IWUSR` - Anyone can read; Modifiable by root
- Can navigate to `/sys/module` to view parameter

/ include / linux / stat.h

```
1  /* SPDX-License-Identifier: GPL-2.0 */
2  #ifndef _LINUX_STAT_H
3  #define _LINUX_STAT_H
4
5
6  #include <asm/stat.h>
7  #include <uapi/linux/stat.h>
8
9  #define S_IRWXUGO      (S_IRWXU|S_IRWXG|S_IRWXO)
10 #define S_IALLUGO      (S_ISUID|S_ISGID|S_ISVTX|S_IRWXUGO)
11 #define S_IRUGO        (S_IRUSR|S_IRGRP|S_IROTH)
12 #define S_IWUGO        (S_IWUSR|S_IWGRP|S_IWOTH)
13 #define S_IXUGO        (S_IXUSR|S_IXGRP|S_IXOTH)
```

Linux Source via Bootlin Elixir Cross Referencer

<https://elixir.bootlin.com/linux/latest/source/include/linux/stat.h>

## Let's Add a Parameter (**multiplier**) Via **module\_param()**

```
#include <linux/module.h>
#include <linux/init.h>
#include <linux/kernel.h>

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Abraham Lincoln");
MODULE_DESCRIPTION("Greatest module in the world!");
MODULE_VERSION("0.000001");

static int multiplier = 10;
module_param(multiplier, int, S_IRUGO);

static int __init hello_init(void){
    printk(KERN_INFO "Oh hi mark\n");
    return 0;
}

static void __exit hello_exit(void){
    printk(KERN_INFO "sad, but still love you\n");
}

module_init(hello_init);
module_exit(hello_exit);
```

## We Need to Print Our Beautiful Parameter (Because We Love it So Much)

```
#include <linux/module.h>
#include <linux/init.h>
#include <linux/kernel.h>

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Abraham Lincoln");
MODULE_DESCRIPTION("Greatest module in the world!");
MODULE_VERSION("0.000001");

static int multiplier = 10;
module_param(multiplier, int, S_IRUGO);

static int __init hello_init(void){
    printk(KERN_INFO "Oh hi mark - I love Lisa %dX more than you do\n", multiplier);
    return 0;
}

static void __exit hello_exit(void){
    printk(KERN_INFO "sad, but still love Lisa %dX more than you\n", multiplier);
}

module_init(hello_init);
module_exit(hello_exit);
```

# Makefile (Unchanged)

```
obj-m+=hello.o
```

```
all:
```

```
    make -C /lib/modules/$(shell uname -r)/build/ M=$(PWD) modules
```

```
clean:
```

```
    make -C /lib/modules/$(shell uname -r)/build/ M=$(PWD) clean
```

# Try it out!

```
sudo insmod hello.ko multiplier=3
```

```
ls /sys/module/
```

```
ls /sys/module/hello/
```

```
ls /sys/module/hello/parameters/
```

```
cat /sys/module/hello/parameters/multiplier
```

```
sudo rmmod hello.ko
```

# Uploaded Example Files

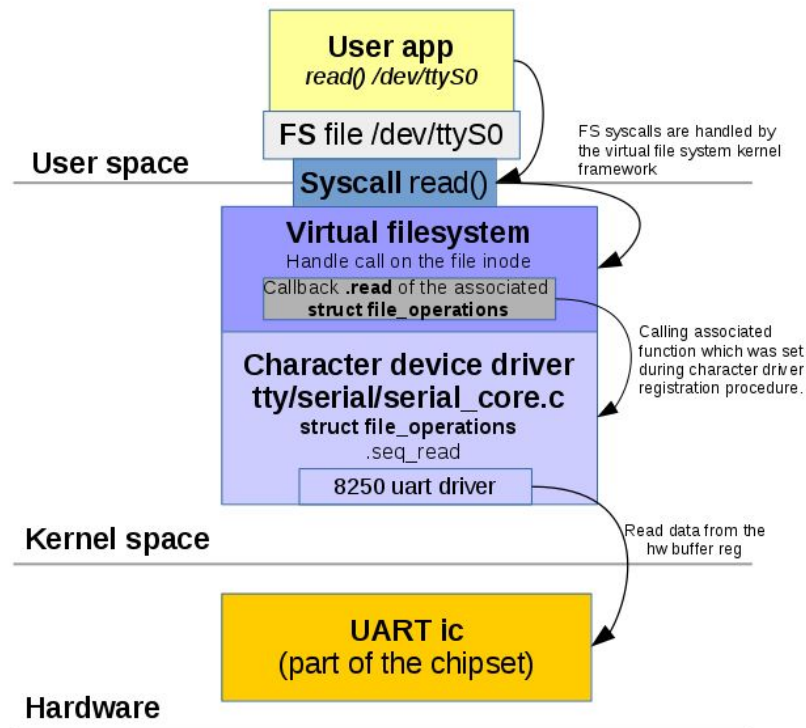
- See Canvas/Examples/Simple\ Module [the \ is intentional]
  - **hello.c**
  - **Makefile**

# Making a Character (Char) Device Driver



# Character (char) Devices (cdevs)

- Interfaced through **stream of bytes**
- System calls commonly used include:
  - open
  - close
  - read
  - write
- Similar to **files**
  - But generally limited to sequential access
- Appear in **/dev**
- Majority of basic devices can be accessed via char drivers



# Demonstration [Notes for Joe]

- Joe goes to Pi
- Inserts module for exercise
- ls /dev
  - Note ttyS0
- ls -al /dev [note c/b/d]
- ls -al ~/hello\_module [note - for regular files]
- Demo expected functionality of module with insmod, ./test, rmmod, and dmesg

# System Calls & Drivers

- User applications can only interact with hardware through system calls to the kernel
- **Drivers implement system calls for devices**
- Our goal today is to write basic code to implement the `open` system call in our driver

# file\_operations

/ include / linux / fs.hAll symbols▼Search Identifier

```
2022 struct file_operations {
2023     struct module *owner;
2024     loff_t (*llseek) (struct file *, loff_t, int);
2025     ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
2026     ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
2027     ssize_t (*read_iter) (struct kiocb *, struct iov_iter *);
2028     ssize_t (*write_iter) (struct kiocb *, struct iov_iter *);
2029     int (*iopoll)(struct kiocb *kiocb, bool spin);
2030     int (*iterate) (struct file *, struct dir_context *);
2031     int (*iterate_shared) (struct file *, struct dir_context *);
2032     __poll_t (*poll) (struct file *,
2033     long (*unlocked_ioctl) (struct file *, unsigned long, long);
2034     long (*compat_ioctl) (struct file *, unsigned long, long);
2035     int (*mmap) (struct file *, struct vm_area_struct *);
2036     unsigned long (*mmap_supported_flags);
2037     int (*open) (struct inode *, struct file *);
2038     int (*flush) (struct file *, fl_owner_t id);
2039     int (*release) (struct inode *, struct file *);
2040     int (*fsync) (struct file *, int, int, int);
2041     int (*fasync) (struct file *, int, int);
2042     int (*fallocate) (struct file *, int, loff_t, loff_t);
```

Function pointer

Exercise 9 Goal: Implement `open` in our driver

# Major Number, Class

- `ls /sys/class/mes`
  - `ls /sys/class/mes/meschar`

# User Interaction

- We need a user application to call **open** for a fake character device we will play with - you will do this in exercise 9
  - Make a C file called testmeschar.c
  - Include the header files stdio.h, stdlib.h, unistd.h, and fcntl.h
  - Inside the main function, print (using printf) a message to the console saying that we are running the program
  - Make a call to the function open()
    - open(), which is part of the standard C library, takes two input arguments
    - The first input is a string that specifies the file path to our fake device we are making, which is at /dev/\*\*\*\* [where \*\*\*\* is the device name you chose]
    - The second argument is O\_RDWR
      - Capital O, not zero; Used to read/write files
    - The output will be an integer; No need to do anything with it
  - return 0

# Open Requires Having The Device At /dev

- Writing a user-space C application that accesses our fake device via our driver through the C `open()` function (which uses the `open` system call)
- Device name will determine what appears at `/dev`
  - E.g., `/dev/COOL_GPU`
- Choose device name (fill in X)
  - `#define DEVICE_NAME "X"`
- Choose class name (fill in X)
  - `#define CLASS_NAME "X"`
- Then can access using `open("/dev/COOL_GPU", O_RDWR)`
  - `O_RDWR` is a macro

# Other Functions In Your Driver

init

- `register_chrdev(0, DEVICE_NAME, &fops)`
- `class_create(THIS_MODULE, CLASS_NAME)`
- `device_create(mescharClass, NULL, MKDEV(majorNumber, 0), NULL, DEVICE_NAME);`

exit

- `device_destroy(mescharClass, MKDEV(majorNumber,0));`
- `class_unregister(mescharClass);`
- `class_destroy(mescharClass);`
- `unregister_chrdev(majorNumber, DEVICE_NAME);`



## Exercise 9: Make new Makefile

obj-m+=hello.o

all:

make -C /lib/modules/\$(shell uname -r)/build/ M=\$(PWD) modules

\$(CC) INSERT\_TEST\_C\_FILE\_NAME\_WITH\_EXTENSION -o test

clean:

make -C /lib/modules/\$(shell uname -r)/build/ M=\$(PWD) clean

rm test