

The background of the slide is a grayscale image of a city street scene, processed using edge detection. The image shows a wide street with a curved path, buildings on the left, and a large, ornate building on the right. The edges are highlighted in white against a black background, creating a high-contrast, abstract representation of the scene.

ELEC 424/553

Mobile & Embedded Systems

Lecture 20
computer vision

Getting Computer Vision (CV) To Work On Your BBB

```
sudo apt update  
sudo apt upgrade
```

```
git clone https://github.com/derekmolloy/boneCV  
git clone https://github.com/jayrambhia/Install-OpenCV.git
```

```
cd ~/Install-OpenCV/Ubuntu/  
./opencv_latest.sh
```

```
sudo apt install libv4l-dev
```

```
cd ~/boneCV  
./build
```

```
sudo apt install ffmpeg
```

```
sudo apt install python3-opencv
```

Follow along at

<https://docs.google.com/document/d/13MDKEBBn4ciiTB9MOuKHkvzUFOXhZ5aDKpfSqapMI9g/e>
[dit#](#)

Housekeeping

- **Project 3 due Friday @ 4pm**
- **The one and only exam will be during the last week of classes**
 - There will be a recorded in-person review session led by me
- **Final project due on the last day of final exams**
 - Team project (3-4 people)
 - Requires Python
 - I will host a help session on Python if requested
 - Graduate students will have extra requirement(s)


Housekeeping

- **Midterm** posted
 - Due Monday November 13th at 11:59pm
- Lecture schedule
 - **Mon Nov 13**: No lecture - I'll be on zoom for any midterm clarification questions
 - **Wed Nov 15**: No lecture - I am going to Denmark!
 - **Mon Nov 20**: Zoom lecture [exercise 14]
 - **Wed Nov 22**: No lecture - *Thanksgiving (food overflow)*
 - **Mon Nov 27**: No lecture - Office hours for final project
 - **Wed Nov 29**: No lecture - Office hours for final project

Housekeeping (continued)

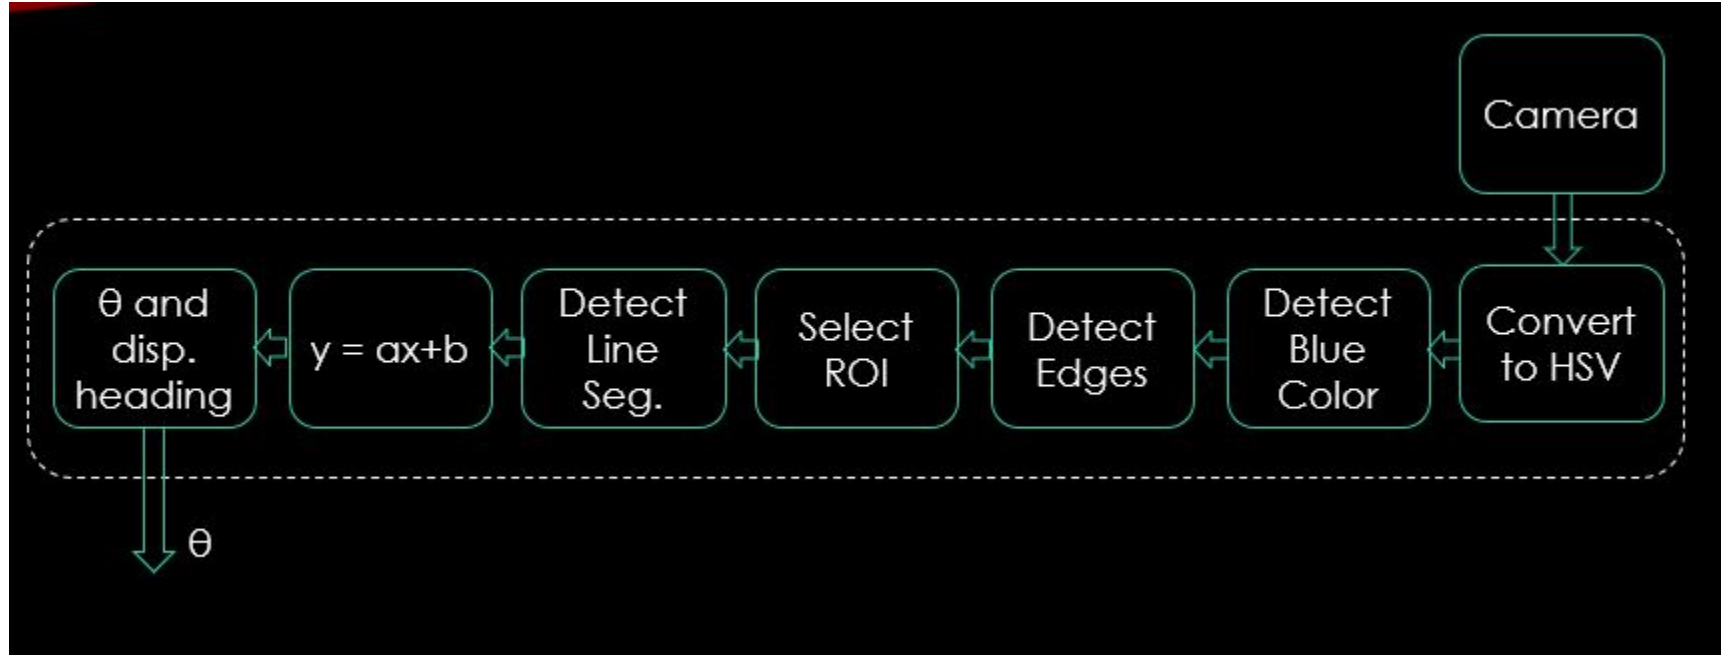
- **Project 3**
 - Teams announced
 - Motor control
 - Goodies will be given
- **Final Project**
 - Autonomous lane-keeping RC car
 - Due Fri Dec 1 (last day of classes)
- **Assignment 3** (optional)
 - Rust
 - Replace assignment or project grade of your choosing (not final project though)
 - TBA on posting

Making Our Own Computer Vision Application



<https://www.instructables.com/Autonomous-Lane-Keeping-Car-Using-Raspberry-Pi-and/>

Pipeline of Image Processing From Instructable







What Do We Call This Approach

- Pipeline approach
- Guess who uses pipeline approaches?
- Tesla
- Likely ever other major player apart from Comma AI

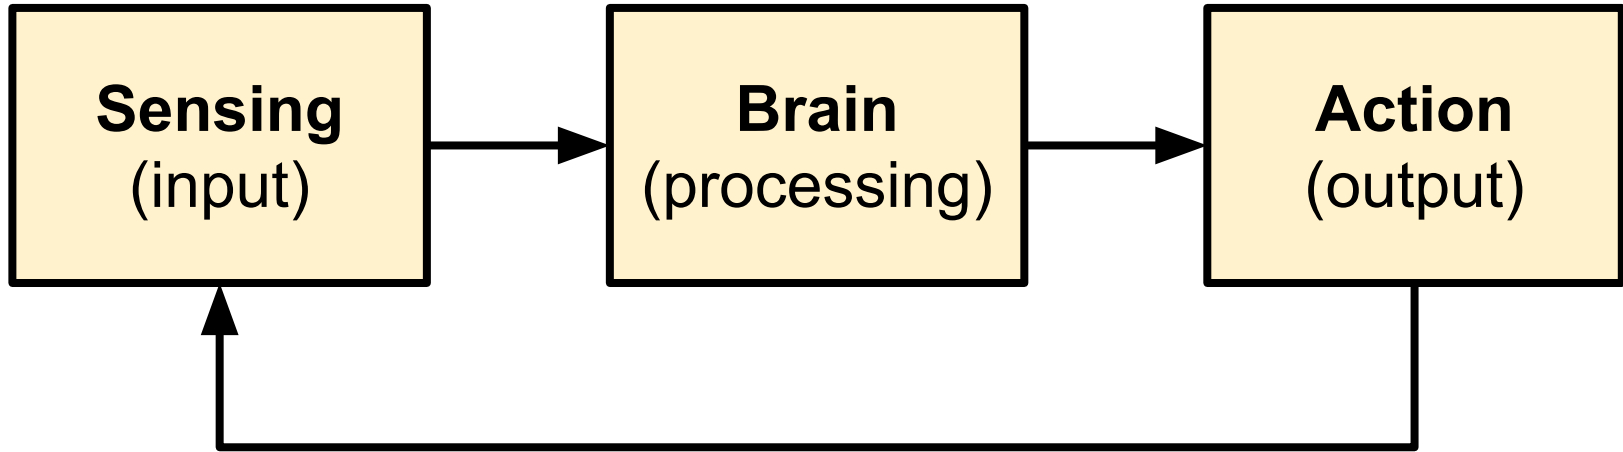
Life - What's It All About?




Humans (Sometimes) Use A Little More Than This



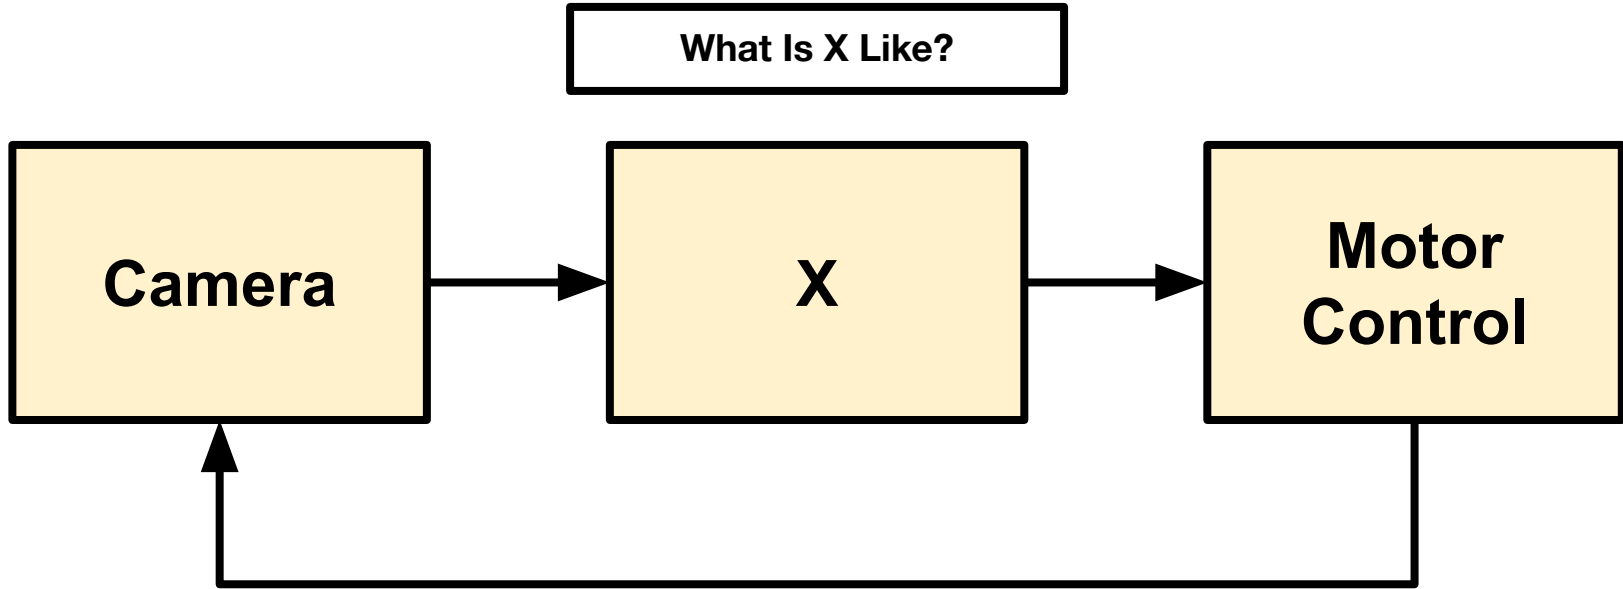
Humans (Sometimes) Use A Little More Than This



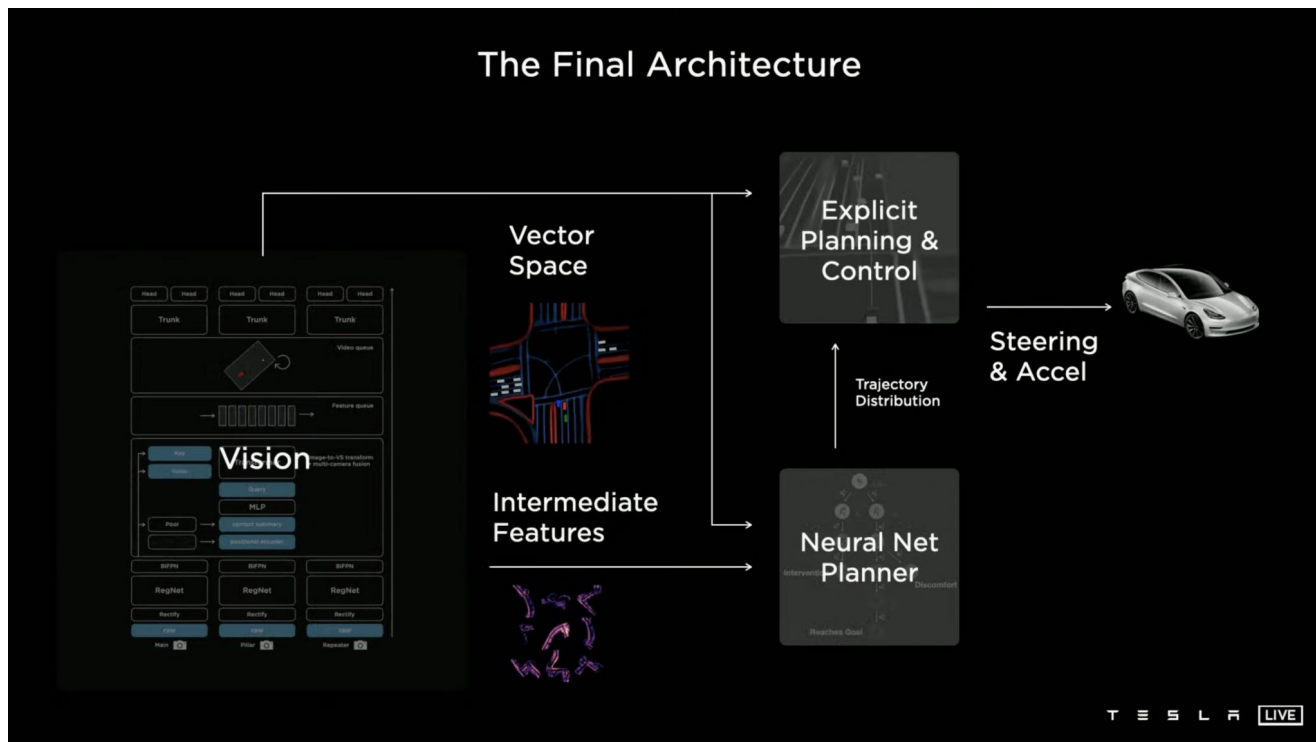
Autonomous Vehicle: What's The Fundamental Goal?




We Need Something More



Tesla's X



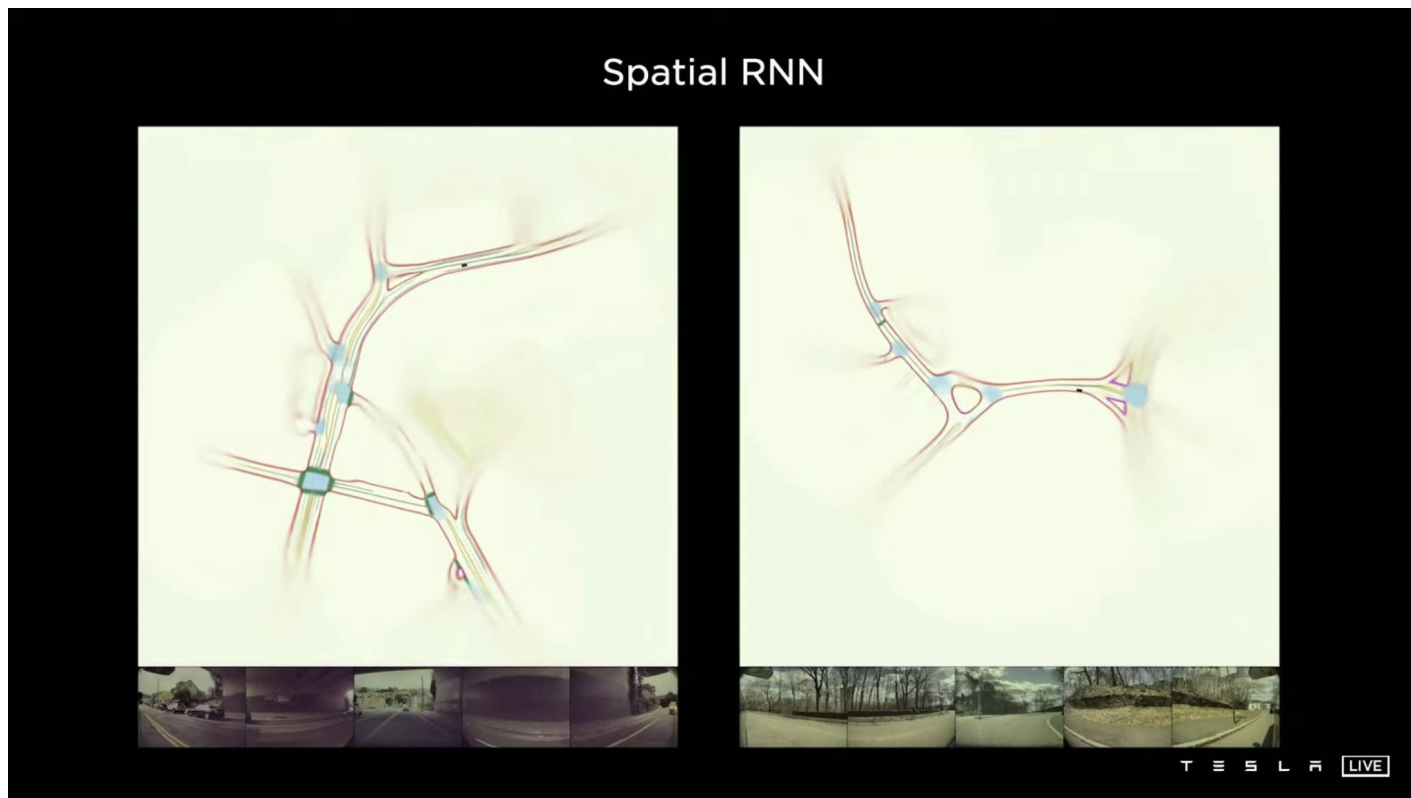
Tesla's X



From Tesla's AI Day event: Teslascope thread on Twitter. URL:


https://twitter.com/teslascope/status/1428524117741871113?ref_src=twsrc%5Etfw%7Ctwcamp%5Etweetembed%7Ctwterm%5E1428524117741871113%7Ctwgr%5E%7Ctwcon%5Es1_&ref_url=https%3A%2F%2Fwww.teslarati.com%2FTesla-ai-day-live-blog%2F

Tesla's X

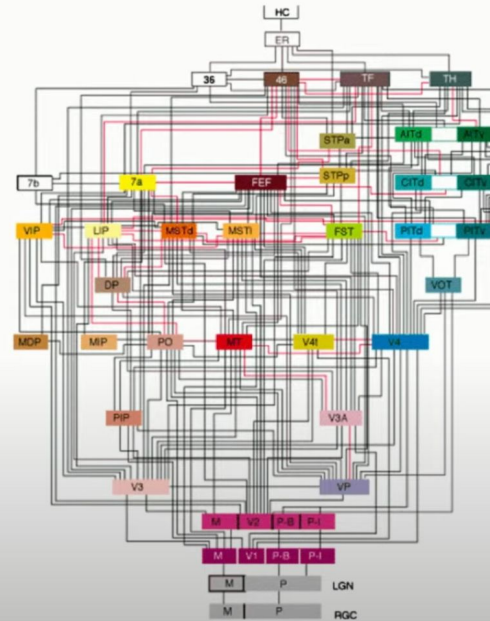



Tesla's X

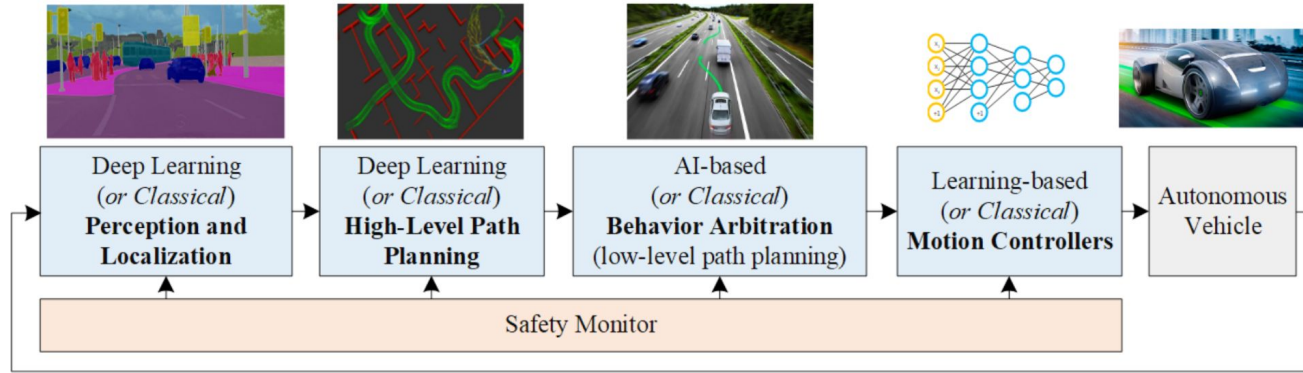
Biological Visual Cortex Wiring



Top-down influences on visual processing. Nature Reviews Neuroscience, 2013.



Tesla's X



Comma AI's X

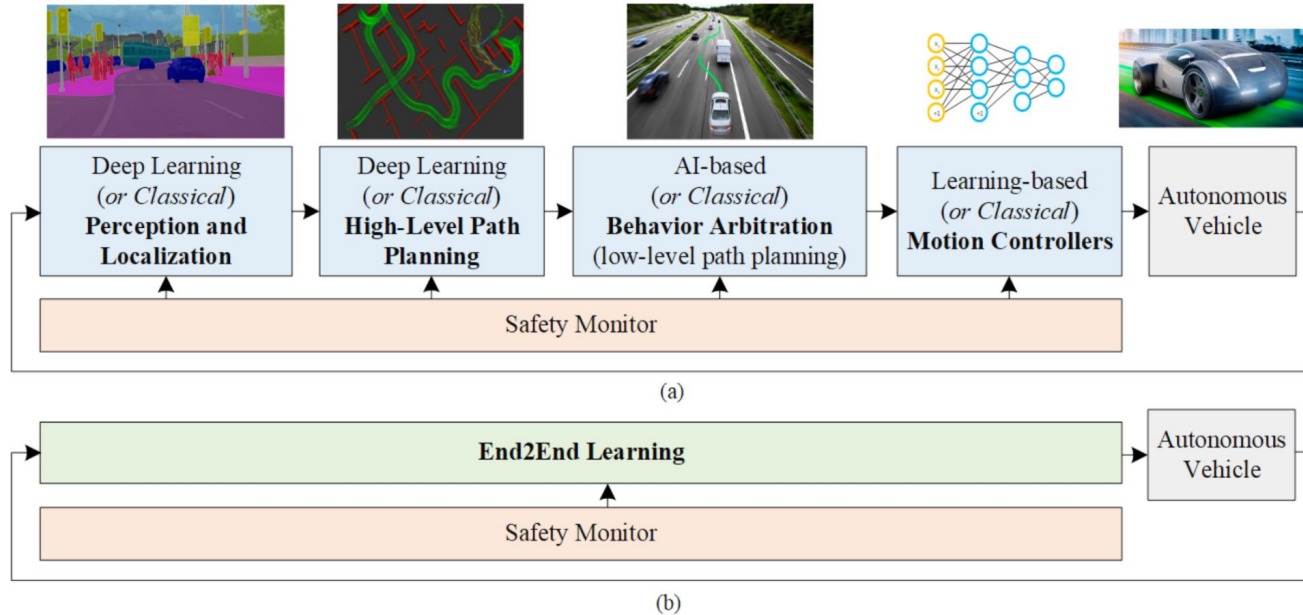



Figure 1: **Deep Learning based self-driving car.** The architecture can be implemented either as a sequential perception-planning-action pipeline (a), or as an End2End system (b). In the sequential pipeline case, the components can be designed either using AI and deep learning methodologies, or based on classical non-learning approaches. End2End learning systems are mainly based on deep learning methods. A safety monitor is usually designed to ensure the safety of each module.



Our X




Camera



HSV





CC Attribution-Share Alike 3.0 Unported license. Unmodified.
https://commons.wikimedia.org/wiki/File:HSV_color_solid_cylinder.png




<https://stackoverflow.com/questions/47483951/how-to-define-a-threshold-value-to-detect-only-green-colour-objects-in-an-image/47483966#47483966>

Blue




CC Attribution-Share Alike 3.0 Unported license. Unmodified.
https://commons.wikimedia.org/wiki/File:HSV_color_solid_cylinder.png




<https://stackoverflow.com/questions/47483951/how-to-define-a-threshold-value-to-detect-only-green-colour-objects-in-an-image/47483966#47483966>

Blue





```
46 ✓ def detect_edges(frame):  
47     # filter for blue lane lines  
48     hsv = cv2.cvtColor(frame, cv2.COLOR_BGR2HSV)  
49     #cv2.imshow("HSV",hsv)  
50     lower_blue = np.array([90, 120, 0], dtype = "uint8")  
51     upper_blue = np.array([150, 255, 255], dtype="uint8")  
52     mask = cv2.inRange(hsv,lower_blue,upper_blue)
```

Blue




```
46 def detect_edges(frame):  
47     # filter for blue lane lines  
48     hsv = cv2.cvtColor(frame, cv2.COLOR_BGR2HSV)  
49     #cv2.imshow("HSV",hsv)  
50     lower_blue = np.array([90, 120, 0], dtype="uint8")  
51     upper_blue = np.array([150, 255, 255], dtype="uint8")  
52     mask = cv2.inRange(hsv, lower_blue, upper_blue)
```




<https://stackoverflow.com/questions/47483951/how-to-define-a-threshold-value-to-detect-only-green-colour-objects-in-an-image/47483966#47483966>


CC Attribution-Share Alike 3.0 Unported license. Unmodified.
https://commons.wikimedia.org/wiki/File:HSV_color_solid_cylinder.png

Edges



Reposition Figure...






The Following Slides On CV Are Inspired By These Lectures

Edge Detection

Convert a 2D image into a set of points where image intensity changes rapidly.

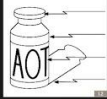
Topics:

- (1) What is an Edge?
- (2) Edge Detection Using Gradients
- (3) Edge Detection Using Laplacian
- (4) Canny Edge Detector
- (5) Corner Detection




Causes of Edges

Rapid changes in image intensity are caused by various physical phenomena.




- Surface Normal Discontinuity
- Depth Discontinuity
- Surface Reflectance Discontinuity
- Illumination Discontinuity






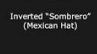
Gradient (∇) as Edge Detector


Gradient Magnitude $S = \|\nabla I\| = \sqrt{\left(\frac{\partial I}{\partial x}\right)^2 + \left(\frac{\partial I}{\partial y}\right)^2}$

Gradient Orientation $\theta = \tan^{-1}\left(\frac{\partial I}{\partial y} / \frac{\partial I}{\partial x}\right)$










Gradient vs. Laplacian

Derivative of Gaussian (∇G)	Laplacian of Gaussian ($\nabla^2 G$)
 $\frac{\partial}{\partial x}(I_{\sigma})$	 $\frac{\partial^2}{\partial x^2}(I_{\sigma}) + \frac{\partial^2}{\partial y^2}(I_{\sigma})$
 $\frac{\partial}{\partial y}(I_{\sigma})$	 Inverted "Sombrero" (Mexican Hat)




Canny Edge Detector Results

Image	$\sigma = 1$
	
	
	

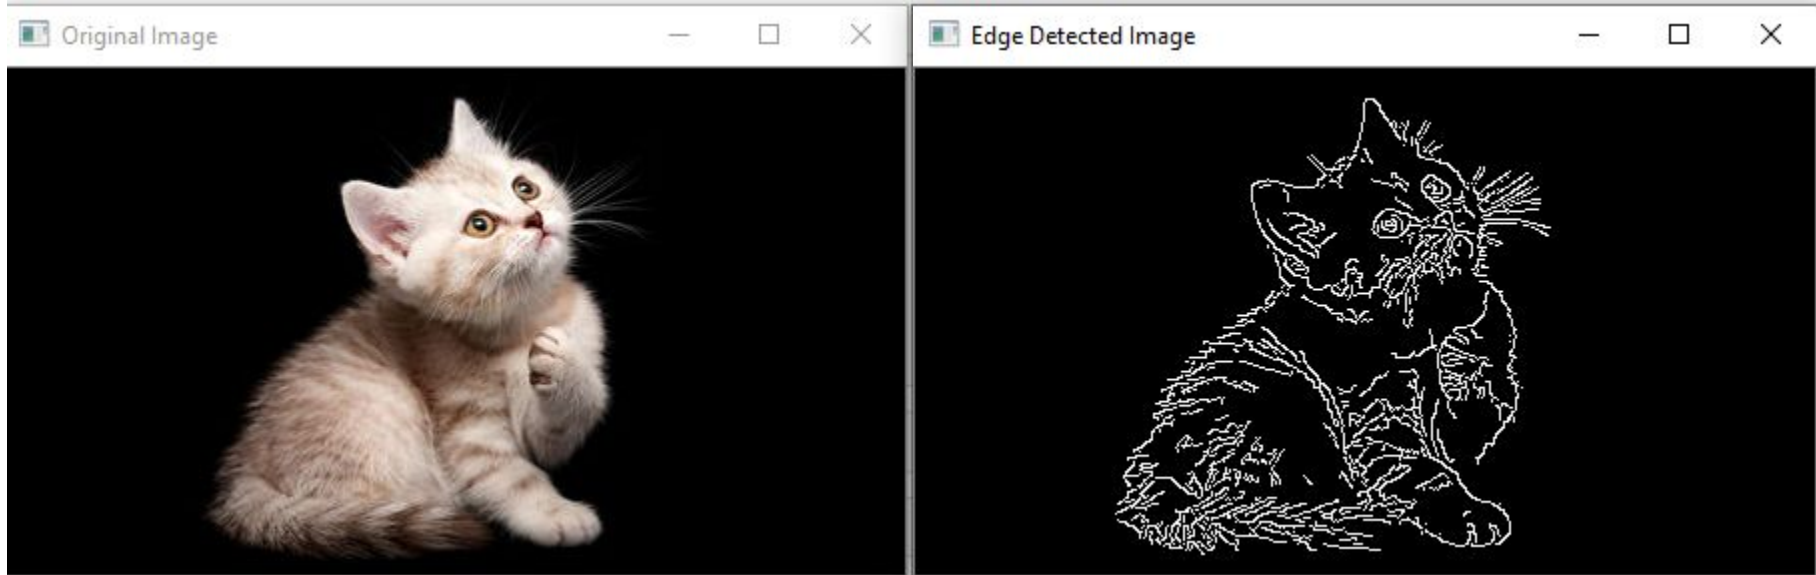



YouTube Channel: [First Principles of Computer Vision](#)




Goal of Edge Detection

What is an edge?






Goal of Edge Detection (2)



Note: OpenCV & waitKey

- You often (always?) have to add `cv2.waitKey(1)` in your loop if you are continually showing images using openCV
- Otherwise you may not get any image window to appear
 - I had this issue on my Mac
- More info - MLK, URL:
<https://machinelearningknowledge.ai/opencv-cv2-waitkey-tutorial-with-examples/>




Edge Detection Specifics

- Often done with black and white (BW) images
- Looking for:
 - Sharp contrast
 - Sharp changes in pixel intensity
 - High frequency activity
- Consider that a BW image is just a matrix of pixel intensities/brightness values


183	190	191	215	195	206	224	166
169	174	139	178	173	150	201	227
172	152	225	216	197	230	196	187
221	176	164	137	153	190	233	192
195	223	200	191	174	217	160	152
207	189	137	189	167	183	224	169
206	206	205	148	234	232	157	150
143	185	144	208	139	182	169	229

8-bit image: Range of 0 to 255

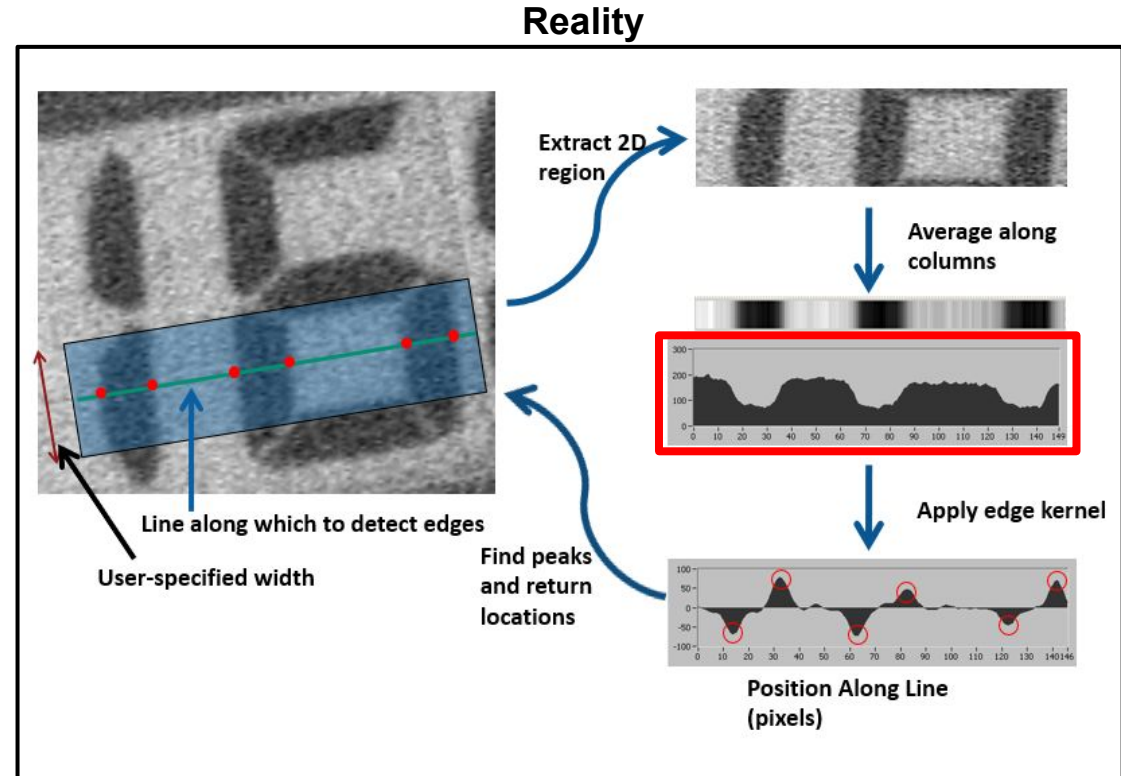
<https://medium.com/@damiandn/an-introduction-to-biological-image-processing-in-imagej-part-1-what-is-an-image-54fc31f3d02d>




What Do Actual Edges Look Like?



“Zero Crossing Detector”. URL:
<https://homepages.inf.ed.ac.uk/rbf/HIPR2/zeros.htm>




“Which Kernel Does the IMAQ Edge Tool 3 Use?”. URL:
<https://knowledge.ni.com/KnowledgeArticleDetails?id=kA00Z0000015CV0SAM&l=en-US>




Gradients & Edge Detection

Edge Detection Using 1st Derivative



Shree K. Nayar

Provides Both Location and Strength of an Edge




Gradients & Edge Detection (2)

Gradient (∇)


Gradient (Partial Derivatives) represents the direction of most rapid change in intensity

$$\nabla I = \left[\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right]$$


Pronounced as “Del I”




$$\nabla I = \left[\frac{\partial I}{\partial x}, 0 \right]$$



$$\nabla I = \left[0, \frac{\partial I}{\partial y} \right]$$



$$\nabla I = \left[\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right]$$





Gradients & Edge Detection (3)

Gradient (∇) as Edge Detector

Gradient Magnitude $S = \|\nabla I\| = \sqrt{\left(\frac{\partial I}{\partial x}\right)^2 + \left(\frac{\partial I}{\partial y}\right)^2}$

Gradient Orientation $\theta = \tan^{-1} \left(\frac{\partial I}{\partial y} / \frac{\partial I}{\partial x} \right)$






Gradients & Edge Detection (4)

Discrete Gradient (∇) Operator

Finite difference approximations:

$$\frac{\partial I}{\partial x} \approx \frac{1}{2\varepsilon} \left((I_{i+1,j+1} - I_{i,j+1}) + (I_{i+1,j} - I_{i,j}) \right)$$

$$\frac{\partial I}{\partial y} \approx \frac{1}{2\varepsilon} \left((I_{i+1,j+1} - I_{i+1,j}) + (I_{i,j+1} - I_{i,j}) \right)$$



Can be implemented as Convolution!


$$\frac{\partial}{\partial x} \approx \frac{1}{2\varepsilon} \begin{bmatrix} -1 & 1 \\ -1 & 1 \end{bmatrix}$$

$$\frac{\partial}{\partial y} \approx \frac{1}{2\varepsilon} \begin{bmatrix} 1 & 1 \\ -1 & -1 \end{bmatrix}$$

Note: Convolution flips have been applied

Shree K. Nayar

Shree K. Nayar



Gradients & Edge Detection (5)

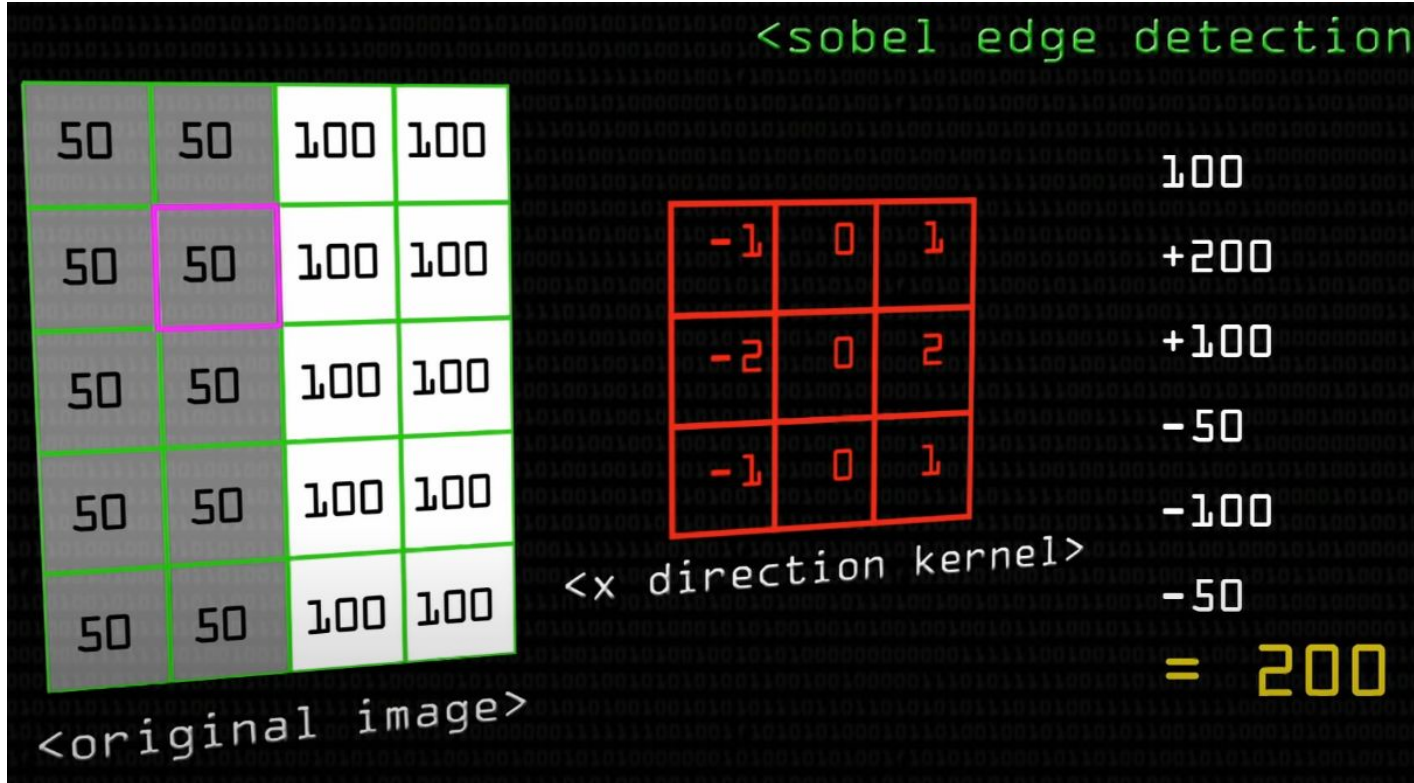




Image would be
BW and
smoothed



Sobel Results





Figure 3: Sobel edge detection




Canny Edge Detection


- Improves upon Sobel
- Sobel feeds into Canny edge detector
- Removes weak edges
- Two steps:
 - Makes edges 1 pixel wide
 - **Hysteresis** - two thresholds





Canny Edge Detection: Hysteresis Thresholding





Sobel Vs. Canny Edge Detection



Figure 3: Sobel edge detection


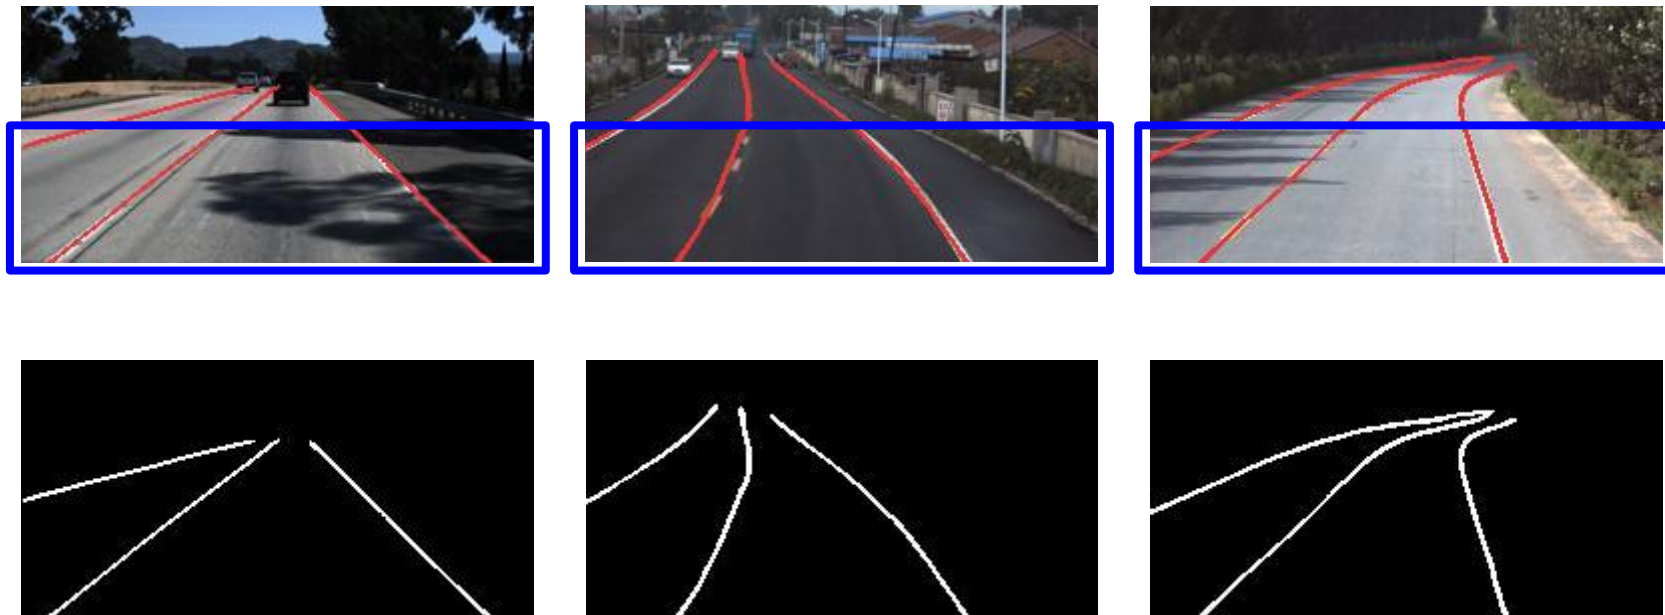



Figure 4: Canny edge detection



ROI: Region of Interest

Do We Need The Whole Scene?





Line Segment Detection

- How do we infer and present a shape?
- We have edges
- But there are holes